

Recurrent Neural Networks

CISC 7026: Introduction to Deep Learning

University of Macau

~~Prof. Qingbiao Li GNN lecture 28 October~~

Admin

~~Prof. Qingbiao Li GNN lecture 28 October~~

November 4 holiday, no lecture

Admin

~~Prof. Qingbiao Li GNN lecture 28 October~~

November 4 holiday, no lecture

Saturday lecture made everyone look sad

Admin

~~Prof. Qingbiao Li GNN lecture 28 October~~

November 4 holiday, no lecture

Saturday lecture made everyone look sad

~~Makeup lecture Saturday October 26, 13:00-16:00~~

Admin

~~Prof. Qingbiao Li GNN lecture 28 October~~

November 4 holiday, no lecture

Saturday lecture made everyone look sad

~~Makeup lecture Saturday October 26, 13:00-16:00~~

Virtual lecture, but department policy states it must be in person

Admin

~~Prof. Qingbiao Li GNN lecture 28 October~~

November 4 holiday, no lecture

Saturday lecture made everyone look sad

~~Makeup lecture Saturday October 26, 13:00-16:00~~

Virtual lecture, but department policy states it must be in person

I will be here on a Saturday (not sure when) to answer questions or discuss the virtual lecture

Admin

~~Prof. Qingbiao Li GNN lecture 28 October~~

November 4 holiday, no lecture

Saturday lecture made everyone look sad

~~Makeup lecture Saturday October 26, 13:00-16:00~~

Virtual lecture, but department policy states it must be in person

I will be here on a Saturday (not sure when) to answer questions or discuss the virtual lecture

Updated assignment 5 is on Moodle, due in 2 weeks

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Agenda


1. **Review**
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Review

In perceptrons, each neuron in a layer is independent

Review

In perceptrons, each neuron in a layer is independent



Review

These images are equivalent to a neural network



Review

These images are equivalent to a neural network



It is a miracle that our neural networks could classify clothing!

Review

A **signal** represents information as a function of time, space or some other variable

Review

A **signal** represents information as a function of time, space or some other variable

$$x(t) = 2t + 1$$

Review

A **signal** represents information as a function of time, space or some other variable

$$x(t) = 2t + 1$$

$$x(u, v) = \frac{u^2}{v} - 3$$

Review

A **signal** represents information as a function of time, space or some other variable

$$x(t) = 2t + 1$$

$$x(u, v) = \frac{u^2}{v} - 3$$

$x(t)$, $x(u, v)$ represent physical processes that we may or may not know

Review

A **signal** represents information as a function of time, space or some other variable

$$x(t) = 2t + 1$$

$$x(u, v) = \frac{u^2}{v} - 3$$

$x(t), x(u, v)$ represent physical processes that we may or may not know


In **signal processing**, we analyze the meaning of signals

Review

$x(t)$ = stock price


Review

$x(t)$ = stock price



Review


$$x(t) = \text{stock price}$$



There is an underlying structure to $x(t)$

Review

$$x(t) = \text{stock price}$$



There is an underlying structure to $x(t)$


Structure: Tomorrow's stock price will be close to today's stock price

Review

$$x(t) = \text{audio}$$


Review

$$x(t) = \text{audio}$$



Review


$$x(t) = \text{audio}$$



Structure: Nearby waves form syllables

Review

$$x(t) = \text{audio}$$



Structure: Nearby waves form syllables

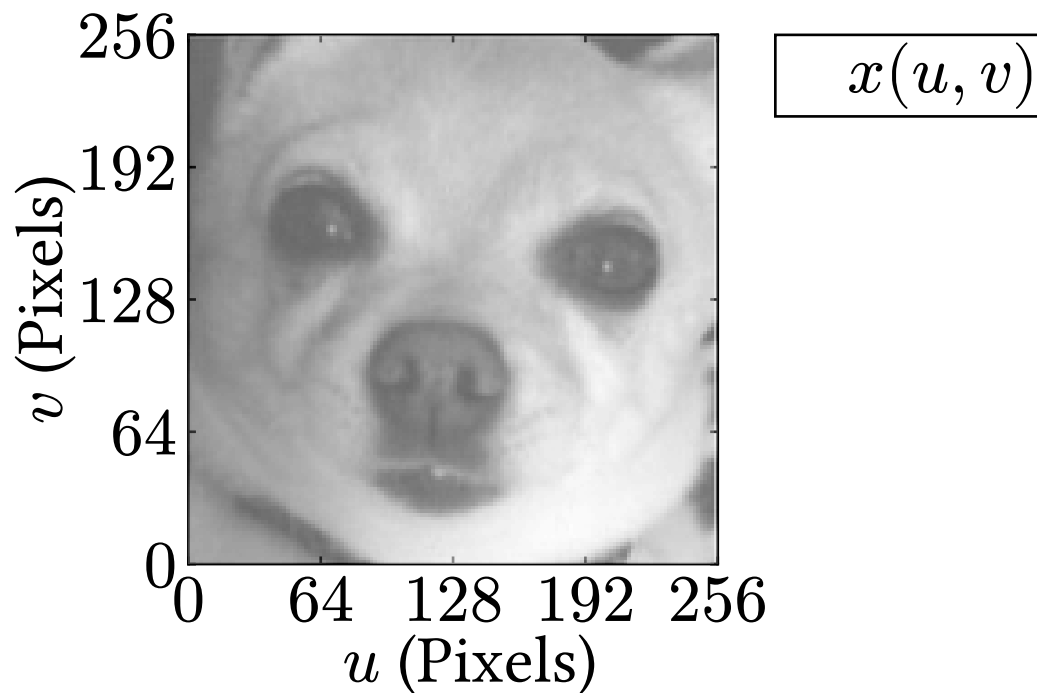
Structure: Nearby syllables combine to create meaning

Review

$$x(u, v) = \text{image}$$


Review

$$x(u, v) = \text{image}$$



Review

$$x(u, v) = \text{image}$$



Structure: Repeated components (circles, symmetry, eyes, nostrils, etc)

Review

Two common properties of signals:

Review

Two common properties of signals:

Locality: Information concentrated over small regions of space/time

Review


Two common properties of signals:

Locality: Information concentrated over small regions of space/time

Translation Equivariance: Shift in signal results in shift in output



Review

A more realistic scenario of locality and translation equivariance





Review

We use convolution to turn signals into useful signals



Review

We use convolution to turn signals into useful signals



Convolution is translation equivariant and local

Review

Convolution is the sum of products of a signal $x(t)$ and a **filter** $g(t)$

Review

Convolution is the sum of products of a signal $x(t)$ and a **filter** $g(t)$

If the t is continuous in $x(t)$

$$x(t) * g(t) = \int_{-\infty}^{\infty} x(t - \tau)g(\tau)d\tau$$

Review

Convolution is the sum of products of a signal $x(t)$ and a **filter** $g(t)$

If the t is continuous in $x(t)$

$$x(t) * g(t) = \int_{-\infty}^{\infty} x(t - \tau)g(\tau)d\tau$$

If the t is discrete in $x(t)$

$$x(t) * g(t) = \sum_{\tau=-\infty}^{\infty} x(t - \tau)g(\tau)$$

Review

Convolution is the sum of products of a signal $x(t)$ and a **filter** $g(t)$

If the t is continuous in $x(t)$


$$x(t) * g(t) = \int_{-\infty}^{\infty} x(t - \tau)g(\tau)d\tau$$

If the t is discrete in $x(t)$


$$x(t) * g(t) = \sum_{\tau=-\infty}^{\infty} x(t - \tau)g(\tau)$$

We slide the filter $g(t)$ across the signal $x(t)$




Review



Review



Review



Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 2 & 1 & & & \\ & & & & \end{bmatrix}$$

Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} \textcolor{red}{1} & \textcolor{red}{2} & 3 & 4 & 5 \\ \textcolor{red}{2} & \textcolor{red}{1} & & & \\ \textcolor{red}{4} & & & & \end{bmatrix}$$

Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ & 2 & 1 & & \\ 4 & 7 & & & \end{bmatrix}$$

Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & \textcolor{red}{3} & \textcolor{red}{4} & 5 \\ & & \textcolor{red}{2} & \textcolor{red}{1} & \\ 4 & 5 & \textcolor{red}{10} & & \end{bmatrix}$$

Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ & & & 2 & 1 \\ 4 & 5 & 10 & 13 & \end{bmatrix}$$

Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 2 & 1 & & & \\ 4 & 5 & 10 & 13 & \end{bmatrix}$$

Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 2 & 1 & & & \\ 4 & 5 & 10 & 13 & \end{bmatrix}$$

To make a convolution layer, we make the filter with trainable parameters

Review

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 2 & 1 & & & \\ 4 & 5 & 10 & 13 & \end{bmatrix}$$

To make a convolution layer, we make the filter with trainable parameters

$$\begin{bmatrix} x(t) \\ g(t) \\ y(t) \end{bmatrix} = \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ \theta_2 & \theta_1 & & & \\ & & & & \end{bmatrix}$$

Review

We can write both a perceptron and convolution in vector form

$$f(x(t), \boldsymbol{\theta}) = \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \\ \vdots \end{bmatrix} \right)$$

Review

We can write both a perceptron and convolution in vector form

$$f(x(t), \boldsymbol{\theta}) = \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \\ \vdots \end{bmatrix} \right) \qquad f(x(t), \boldsymbol{\theta}) = \begin{bmatrix} \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \end{bmatrix} \right) \\ \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.2) \\ x(0.3) \\ \vdots \end{bmatrix} \right) \end{bmatrix}$$

A convolution layer applies a “mini” perceptron to every few timesteps

Review

We can write both a perceptron and convolution in vector form

$$f(x(t), \boldsymbol{\theta}) = \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \\ \vdots \end{bmatrix} \right) \quad f(x(t), \boldsymbol{\theta}) = \begin{bmatrix} \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \end{bmatrix} \right) \\ \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.2) \\ x(0.3) \\ \vdots \end{bmatrix} \right) \end{bmatrix}$$

A convolution layer applies a “mini” perceptron to every few timesteps

The output size depends on the signal length

Review

If we want a single output, we should **pool**

Review

If we want a single output, we should **pool**

$$z(t) = f(x(t), \boldsymbol{\theta}) = \left[\sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x^{(0.1)} \\ x^{(0.2)} \end{bmatrix} \right) \quad \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x^{(0.2)} \\ x^{(0.3)} \end{bmatrix} \right) \quad \dots \right]^\top$$

Review

If we want a single output, we should **pool**

$$z(t) = f(x(t), \boldsymbol{\theta}) = \left[\sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \end{bmatrix} \right) \quad \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.2) \\ x(0.3) \end{bmatrix} \right) \quad \dots \right]^\top$$

$$\text{SumPool}(z(t)) = \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \end{bmatrix} \right) + \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.2) \\ x(0.3) \end{bmatrix} \right) + \dots$$

Review

If we want a single output, we should **pool**

$$z(t) = f(x(t), \boldsymbol{\theta}) = \left[\sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \end{bmatrix} \right) \quad \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.2) \\ x(0.3) \end{bmatrix} \right) \quad \dots \right]^\top$$

$$\text{SumPool}(z(t)) = \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.1) \\ x(0.2) \end{bmatrix} \right) + \sigma \left(\boldsymbol{\theta}^\top \begin{bmatrix} 1 \\ x(0.2) \\ x(0.3) \end{bmatrix} \right) + \dots$$

$$\text{MeanPool}(z(t)) = \frac{1}{T - k + 1} \text{SumPool}(z(t)); \quad \text{MaxPool}(z(t)) = \max(z(t))$$

Review

Our examples considered:

Review

Our examples considered:

- 1 dimensional variable t

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$
- 1 filter

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$
- 1 filter

We can expand to arbitrary dimensions

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$
- 1 filter

We can expand to arbitrary dimensions

- Images 2D (u, v)

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$
- 1 filter

We can expand to arbitrary dimensions

- Images 2D (u, v)
- Video 3D (u, v, t)

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$
- 1 filter

We can expand to arbitrary dimensions

- Images 2D (u, v)
- Video 3D (u, v, t)
- Robot position, orientation, and time (7D)

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$
- 1 filter

We can expand to arbitrary dimensions

- Images 2D (u, v)
- Video 3D (u, v, t)
- Robot position, orientation, and time (7D)
- Arbitrary signals

Review

Our examples considered:

- 1 dimensional variable t
- 1 dimensional output/channel $x(t)$
- 1 filter

We can expand to arbitrary dimensions

- Images 2D (u, v)
- Video 3D (u, v, t)
- Robot position, orientation, and time (7D)
- Arbitrary signals

The idea is exactly the same

Review

0	1	1	1	1	1	0	1
0	1	1	1	1	1	1	1
1	0	0	1	1	0	0	1
1	0	0	1	1	0	0	1
1	1	0	0	0	1	1	1
1	1	0	0	0	1	1	1
0	1	0	0	0	1	1	0
1	0	1	1	1	1	1	1

0	1	1	1	1	1	0	1
0	1	1	1	1	1	1	1
1	0	0	1	1	0	0	1
1	0	0	1	1	0	0	1
1	1	0	0	0	1	1	1
1	1	0	0	0	1	1	1
0	1	0	0	0	1	1	0
1	0	1	1	1	1	1	1

0	1	1	1	1	1	0	1
0	1	1	1	1	1	1	1
1	0	0	1	1	0	0	1
1	0	0	1	1	0	0	1
1	1	0	0	0	1	1	1
1	1	0	0	0	1	1	1
0	1	0	0	0	1	1	0
1	0	1	1	1	1	1	1

$$\begin{array}{|c|c|} \hline 2 & 0 \\ \hline 0 & 1 \\ \hline \end{array}
 \quad + \quad
 \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 0 & 1 \\ \hline \end{array}
 \quad + \quad
 \begin{array}{|c|c|} \hline 1 & 0 \\ \hline 0 & 0 \\ \hline \end{array}$$

Review

One last thing, **stride** allows you to “skip” cells during convolution

Review

One last thing, **stride** allows you to “skip” cells during convolution

This can decrease the size of image without pooling

Review

One last thing, **stride** allows you to “skip” cells during convolution

This can decrease the size of image without pooling

Padding adds zero pixels to the image to increase the output size

0	1	1	1	1	1	0	1
0	1	1	1	1	1	1	1
1	0	0	1	1	0	0	1
1	0	0	1	1	0	0	1
1	1	0	0	0	1	1	1
1	1	0	0	0	1	1	1
0	1	0	0	0	1	1	0
1	0	1	1	1	1	1	1

*

1	0
0	1

Agenda

1. **Review**
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Agenda

1. Review
2. **Sequence Modeling**
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Sequence Modeling

We previously used convolution to model signals

Sequence Modeling

We previously used convolution to model signals

Convolution is an electrical engineering approach to modeling sequences

Sequence Modeling

We previously used convolution to model signals

Convolution is an electrical engineering approach to modeling sequences

Now, we will discuss a neuroscience approach to sequence modeling

Sequence Modeling

We previously used convolution to model signals

Convolution is an electrical engineering approach to modeling sequences

Now, we will discuss a neuroscience approach to sequence modeling

We call these models **recurrent models**

Sequence Modeling

We previously used convolution to model signals

Convolution is an electrical engineering approach to modeling sequences

Now, we will discuss a neuroscience approach to sequence modeling

We call these models **recurrent models**

You can solve temporal tasks using either convolution or RNNs

Sequence Modeling

We previously used convolution to model signals

Convolution is an electrical engineering approach to modeling sequences

Now, we will discuss a neuroscience approach to sequence modeling

We call these models **recurrent models**

You can solve temporal tasks using either convolution or RNNs

So what is the difference between convolution and recurrent models?

Sequence Modeling

Convolution works over inputs of any variables (time, space, etc)

Sequence Modeling

Convolution works over inputs of any variables (time, space, etc)

Recurrent neural networks only work with time

Sequence Modeling

Convolution works over inputs of any variables (time, space, etc)

Recurrent neural networks only work with time

Convolution makes use of locality and translation equivariance properties

Sequence Modeling

Convolution works over inputs of any variables (time, space, etc)

Recurrent neural networks only work with time

Convolution makes use of locality and translation equivariance properties

Recurrent models do not assume locality or equivariance

Sequence Modeling

Convolution works over inputs of any variables (time, space, etc)

Recurrent neural networks only work with time

Convolution makes use of locality and translation equivariance properties

Recurrent models do not assume locality or equivariance

Equivariance and locality make learning more efficient, but not all problems have this structure

Sequence Modeling

Convolution works over inputs of any variables (time, space, etc)

Recurrent neural networks only work with time

Convolution makes use of locality and translation equivariance properties

Recurrent models do not assume locality or equivariance

Equivariance and locality make learning more efficient, but not all problems have this structure

Let us examine some real life signals, and see if these properties hold

Sequence Modeling

Example 1: You like dinosaurs as a child, you grow up and study dinosaurs for work

Sequence Modeling

Example 1: You like dinosaurs as a child, you grow up and study dinosaurs for work

Question: Is this local?

Sequence Modeling

Example 1: You like dinosaurs as a child, you grow up and study dinosaurs for work

Question: Is this local?

Answer: No, two related events separated by 20 years

Sequence Modeling

Example 2: Your parent changes your diaper

Sequence Modeling

Example 2: Your parent changes your diaper

Question: Translation equivariant?

Sequence Modeling

Example 2: Your parent changes your diaper

Question: Translation equivariant?

No! Ok if you are a baby, different meaning if you are an adult!

Sequence Modeling

Example 3: You hear a gunshot
then see runners



Sequence Modeling


Example 3: You hear a gunshot
then see runners



Question: Translation equivariant?

Sequence Modeling

Example 3: You hear a gunshot
then see runners



Question: Translation equivariant?

Answer: No! (1) gunshot, (2) see runners, enjoy the race. (1) see runners, (2) hear gunshot, you start running too!

Sequence Modeling

Example 3: You hear a gunshot
then see runners



Question: Translation equivariant?

Answer: No! (1) gunshot, (2) see runners, enjoy the race. (1) see runners, (2) hear gunshot, you start running too!

Question: Any other examples?

Sequence Modeling

Problems without locality and translation equivariance are difficult to solve with convolution

Sequence Modeling

Problems without locality and translation equivariance are difficult to solve with convolution

For these problems, we need something else!

Sequence Modeling

Problems without locality and translation equivariance are difficult to solve with convolution

For these problems, we need something else!

How do humans experience time and process temporal data?

Sequence Modeling

Problems without locality and translation equivariance are difficult to solve with convolution

For these problems, we need something else!

How do humans experience time and process temporal data?

Can we design a neural network based on human perceptions of time?

Agenda

1. Review
2. **Sequence Modeling**
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Agenda

1. Review
2. Sequence Modeling
3. **Composite Memory**
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Composite Memory

How do humans process temporal data?

Composite Memory

How do humans process temporal data?

We only perceive the present

Composite Memory

How do humans process temporal data?

We only perceive the present

See dog → photoreceptors fire → neurons fire in the brain

Composite Memory

How do humans process temporal data?

We only perceive the present

See dog \rightarrow photoreceptors fire \rightarrow neurons fire in the brain

No dog \rightarrow no photoreceptors fire \rightarrow no neurons fire

Composite Memory

How do humans process temporal data?

We only perceive the present

See dog \rightarrow photoreceptors fire \rightarrow neurons fire in the brain

No dog \rightarrow no photoreceptors fire \rightarrow no neurons fire

We know there was a dog, even if we no longer see it

Composite Memory

How do humans process temporal data?

We only perceive the present

See dog \rightarrow photoreceptors fire \rightarrow neurons fire in the brain

No dog \rightarrow no photoreceptors fire \rightarrow no neurons fire

We know there was a dog, even if we no longer see it

We can reason over time by recording information as **memories**

Composite Memory

How do humans process temporal data?

We only perceive the present

See dog \rightarrow photoreceptors fire \rightarrow neurons fire in the brain

No dog \rightarrow no photoreceptors fire \rightarrow no neurons fire

We know there was a dog, even if we no longer see it

We can reason over time by recording information as **memories**

Humans process temporal data by storing and recalling memories

Composite Memory



John Locke (1690) believed that
consciousness and identity arise
from memories

Composite Memory



John Locke (1690) believed that
consciousness and identity arise
from memories

If all your memories were erased,
you would be a different person

Composite Memory



John Locke (1690) believed that
consciousness and identity arise
from memories

If all your memories were erased,
you would be a different person


Without the ability to reason over
memories, we would only react to
stimuli like bacteria

Composite Memory

So how do we model memory in humans?


Composite Memory

Francis Galton (1822-1911)
photo composite memory




Composite Memory

Francis Galton (1822-1911)
photo composite memory



Composite photo of members of a
party



Composite Memory

Task: Find a mathematical model of how our mind represents memories

Composite Memory

Task: Find a mathematical model of how our mind represents memories

$X : \mathbb{R}^{h \times w}$ People you see at the party

Composite Memory

Task: Find a mathematical model of how our mind represents memories

$X : \mathbb{R}^{h \times w}$ People you see at the party

$H : \mathbb{R}^{h \times w}$ The image in your mind

$$f : X^T \times \Theta \mapsto H$$

Composite Memory

Task: Find a mathematical model of how our mind represents memories

$X : \mathbb{R}^{h \times w}$ People you see at the party

$H : \mathbb{R}^{h \times w}$ The image in your mind

$$f : X^T \times \Theta \mapsto H$$

Composite photography/memory uses a weighted sum

Composite Memory

Task: Find a mathematical model of how our mind represents memories

$X : \mathbb{R}^{h \times w}$ People you see at the party

$H : \mathbb{R}^{h \times w}$ The image in your mind

$$f : X^T \times \Theta \mapsto H$$

Composite photography/memory uses a weighted sum

$$f(x, \theta) = \sum_{i=1}^T \theta^\top \bar{x}_i$$

Composite Memory

$$f(\boldsymbol{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \overline{\boldsymbol{x}}_i$$

Composite Memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \overline{\mathbf{x}}_i$$

What if we see a new face?

Composite Memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

What if we see a new face?

$$f(\mathbf{x}, \boldsymbol{\theta}) = \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}}$$

Composite Memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

What if we see a new face?

$$f(\mathbf{x}, \boldsymbol{\theta}) = \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}}$$

And another new face?

Composite Memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

What if we see a new face?

$$f(\mathbf{x}, \boldsymbol{\theta}) = \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}}$$

And another new face?

$$= \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{newnew}}$$

Composite Memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

What if we see a new face?

$$f(\mathbf{x}, \boldsymbol{\theta}) = \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}}$$

And another new face?

$$= \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{newnew}}$$

We repeat the same process for each new face

Composite Memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

What if we see a new face?

$$f(\mathbf{x}, \boldsymbol{\theta}) = \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}}$$

And another new face?

$$= \left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i \right) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{new}} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{\text{newnew}}$$

We repeat the same process for each new face

We can rewrite f as a **recurrent function**

Agenda

1. Review
2. Sequence Modeling
3. **Composite Memory**
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. **Linear Recurrence**
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Linear Recurrence

Let us rewrite composite memory as a recurrent function

Linear Recurrence

Let us rewrite composite memory as a recurrent function

$$f(\boldsymbol{x}, \boldsymbol{\theta}) = \underbrace{\left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_i \right)}_h + \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_{\text{new}}$$

Linear Recurrence

Let us rewrite composite memory as a recurrent function

$$f(\boldsymbol{x}, \boldsymbol{\theta}) = \underbrace{\left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_i \right)}_h + \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_{\text{new}}$$

$$f(h, \boldsymbol{x}, \boldsymbol{\theta}) = h + \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}$$

Linear Recurrence

Let us rewrite composite memory as a recurrent function

$$f(\boldsymbol{x}, \boldsymbol{\theta}) = \underbrace{\left(\sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_i \right)}_{\boldsymbol{h}} + \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_{\text{new}}$$

$$f(\boldsymbol{h}, \boldsymbol{x}, \boldsymbol{\theta}) = \boldsymbol{h} + \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}$$

$$\boldsymbol{x} \in \mathbb{R}^{d_x}, \quad \boldsymbol{h} \in \mathbb{R}^{d_h}$$

Linear Recurrence

$$\boldsymbol{x} \in \mathbb{R}^{d_x}, \quad \boldsymbol{h} \in \mathbb{R}^{d_h}$$

$$f(\boldsymbol{h}, \boldsymbol{x}, \boldsymbol{\theta}) = \boldsymbol{h} + \boldsymbol{\theta}^\top \overline{\boldsymbol{x}}$$

Linear Recurrence

$$\boldsymbol{x} \in \mathbb{R}^{d_x}, \quad \boldsymbol{h} \in \mathbb{R}^{d_h}$$

$$f(\boldsymbol{h}, \boldsymbol{x}, \boldsymbol{\theta}) = \boldsymbol{h} + \boldsymbol{\theta}^\top \overline{\boldsymbol{x}}$$

$$\textcolor{red}{h}_1 = f(\mathbf{0}, \boldsymbol{x}_1, \boldsymbol{\theta}) = \mathbf{0} + \boldsymbol{\theta}^\top \overline{\boldsymbol{x}}_1$$

Linear Recurrence

$$\mathbf{x} \in \mathbb{R}^{d_x}, \quad \mathbf{h} \in \mathbb{R}^{d_h}$$

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}$$

$$\mathbf{h}_1 = f(\mathbf{0}, \mathbf{x}_1, \boldsymbol{\theta}) = \mathbf{0} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_1$$

$$\mathbf{h}_2 = f(\mathbf{h}_1, \mathbf{x}_2, \boldsymbol{\theta}) = \mathbf{h}_1 + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_2$$

Linear Recurrence

$$\mathbf{x} \in \mathbb{R}^{d_x}, \quad \mathbf{h} \in \mathbb{R}^{d_h}$$

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}$$

$$\mathbf{h}_1 = f(\mathbf{0}, \mathbf{x}_1, \boldsymbol{\theta}) = \mathbf{0} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_1$$

$$\mathbf{h}_2 = f(\mathbf{h}_1, \mathbf{x}_2, \boldsymbol{\theta}) = \mathbf{h}_1 + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_2$$

$$\mathbf{h}_3 = f(\mathbf{h}_2, \mathbf{x}_3, \boldsymbol{\theta}) = \mathbf{h}_2 + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_3$$

Linear Recurrence

$$\mathbf{x} \in \mathbb{R}^{d_x}, \quad \mathbf{h} \in \mathbb{R}^{d_h}$$

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}$$

$$\textcolor{red}{h}_1 = f(\mathbf{0}, \mathbf{x}_1, \boldsymbol{\theta}) = \mathbf{0} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_1$$

$$\textcolor{green}{h}_2 = f(\textcolor{red}{h}_1, \mathbf{x}_2, \boldsymbol{\theta}) = h_1 + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_2$$

$$h_3 = f(\textcolor{green}{h}_2, \mathbf{x}_3, \boldsymbol{\theta}) = h_2 + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_3$$

\vdots

$$\mathbf{y} = h_T = f(h_{T-1}, \mathbf{x}_T, \boldsymbol{\theta}) = h_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence

$$\mathbf{x} \in \mathbb{R}^{d_x}, \quad \mathbf{h} \in \mathbb{R}^{d_h}$$

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}$$

$$\textcolor{red}{h}_1 = f(\mathbf{0}, \mathbf{x}_1, \boldsymbol{\theta}) = \mathbf{0} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_1$$

$$\textcolor{green}{h}_2 = f(\textcolor{red}{h}_1, \mathbf{x}_2, \boldsymbol{\theta}) = h_1 + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_2$$

$$h_3 = f(\textcolor{green}{h}_2, \mathbf{x}_3, \boldsymbol{\theta}) = h_2 + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_3$$

\vdots

$$\mathbf{y} = h_T = f(h_{T-1}, \mathbf{x}_T, \boldsymbol{\theta}) = h_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence

Question: What is the meaning of h in humans?

Linear Recurrence

Question: What is the meaning of h in humans?



Linear Recurrence

Right now, our model remembers everything

Linear Recurrence

Right now, our model remembers everything

But h is a fixed size, what if T is very large?

Linear Recurrence

Right now, our model remembers everything

But h is a fixed size, what if T is very large?

If we keep adding and adding x into h , we will run out of space

<https://www.youtube.com/watch?v=IQ8Aak-k5Yc>

Linear Recurrence

Right now, our model remembers everything

But h is a fixed size, what if T is very large?

If we keep adding and adding x into h , we will run out of space

<https://www.youtube.com/watch?v=IQ8Aak-k5Yc>

Humans cannot remember everything!

Linear Recurrence

Right now, our model remembers everything

But h is a fixed size, what if T is very large?

If we keep adding and adding x into h , we will run out of space


<https://www.youtube.com/watch?v=IQ8Aak-k5Yc>

Humans cannot remember everything!

We forget old information

Linear Recurrence


Murdock (1982)



Linear Recurrence


Murdock (1982)

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$



Linear Recurrence

Murdock (1982)




$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

Key Idea: $\lim_{T \rightarrow \infty} \gamma^T = 0$

Linear Recurrence

Murdock (1982)




$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

Key Idea: $\lim_{T \rightarrow \infty} \gamma^T = 0$

Let $\gamma = 0.9$

Linear Recurrence

Murdock (1982)



$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$


Key Idea: $\lim_{T \rightarrow \infty} \gamma^T = 0$

Let $\gamma = 0.9$

$$0.9 \cdot 0.9 \cdot \mathbf{h} = 0.81\mathbf{h}$$

Linear Recurrence

Murdock (1982)



$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

Key Idea: $\lim_{T \rightarrow \infty} \gamma^T = 0$


Let $\gamma = 0.9$

$$0.9 \cdot 0.9 \cdot \mathbf{h} = 0.81\mathbf{h}$$

$$0.9 \cdot 0.9 \cdot 0.9 \cdot \mathbf{h} = 0.729\mathbf{h}$$

Linear Recurrence

Murdock (1982)



$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

Key Idea: $\lim_{T \rightarrow \infty} \gamma^T = 0$

Let $\gamma = 0.9$


$$0.9 \cdot 0.9 \cdot \mathbf{h} = 0.81\mathbf{h}$$

$$0.9 \cdot 0.9 \cdot 0.9 \cdot \mathbf{h} = 0.729\mathbf{h}$$

$$0.9 \cdot 0.9 \cdot 0.9 \cdot \dots \cdot \mathbf{h} = 0$$

Linear Recurrence

Murdock (1982)



$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

Key Idea: $\lim_{T \rightarrow \infty} \gamma^T = 0$

Let $\gamma = 0.9$

$$0.9 \cdot 0.9 \cdot \mathbf{h} = 0.81\mathbf{h}$$

$$0.9 \cdot 0.9 \cdot 0.9 \cdot \mathbf{h} = 0.729\mathbf{h}$$

$$0.9 \cdot 0.9 \cdot 0.9 \cdot \dots \cdot \mathbf{h} = 0$$

Let us work out how forgetting works

Linear Recurrence

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

Linear Recurrence

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

$$\mathbf{h}_T = \gamma \mathbf{h}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

$$\mathbf{h}_T = \gamma \mathbf{h}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Question: What is \mathbf{h}_{T-1} in terms of \mathbf{h}_{T-2} ?

Linear Recurrence

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

$$\mathbf{h}_T = \gamma \mathbf{h}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Question: What is \mathbf{h}_{T-1} in terms of \mathbf{h}_{T-2} ?

$$\mathbf{h}_{T-1} = \gamma \mathbf{h}_{T-2} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}$$

Linear Recurrence

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

$$\mathbf{h}_T = \gamma \mathbf{h}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Question: What is \mathbf{h}_{T-1} in terms of \mathbf{h}_{T-2} ?

$$\mathbf{h}_{T-1} = \gamma \mathbf{h}_{T-2} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}$$

$$\mathbf{h}_T = \gamma(\gamma \mathbf{h}_{T-2} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

$$\mathbf{h}_T = \gamma \mathbf{h}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Question: What is \mathbf{h}_{T-1} in terms of \mathbf{h}_{T-2} ?

$$\mathbf{h}_{T-1} = \gamma \mathbf{h}_{T-2} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}$$

$$\mathbf{h}_T = \gamma(\gamma \mathbf{h}_{T-2} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

$$\mathbf{h}_T = \gamma(\gamma(\gamma \mathbf{h}_{T-3} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2}) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}; \quad 0 < \gamma < 1$$

$$\mathbf{h}_T = \gamma \mathbf{h}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Question: What is \mathbf{h}_{T-1} in terms of \mathbf{h}_{T-2} ?

$$\mathbf{h}_{T-1} = \gamma \mathbf{h}_{T-2} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}$$

$$\mathbf{h}_T = \gamma(\gamma \mathbf{h}_{T-2} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

$$\mathbf{h}_T = \gamma(\gamma(\gamma \mathbf{h}_{T-3} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2}) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1}) + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

$$\mathbf{h}_T = \gamma^3 \mathbf{h}_{T-3} + \gamma^2 \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2} + \gamma \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence

$$\mathbf{h}_T = \gamma^3 \mathbf{h}_{T-3} + \gamma^2 \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2} + \gamma \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence


$$\mathbf{h}_T = \gamma^3 \mathbf{h}_{T-3} + \gamma^2 \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2} + \gamma \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

$$\mathbf{h}_T = \gamma^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_1 + \gamma^{T-1} \boldsymbol{\theta}^\top \bar{\mathbf{x}}_2 + \dots + \gamma^2 \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2} + \gamma \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

Linear Recurrence

$$\mathbf{h}_T = \gamma^3 \mathbf{h}_{T-3} + \gamma^2 \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2} + \gamma \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$

$$\mathbf{h}_T = \gamma^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_1 + \gamma^{T-1} \boldsymbol{\theta}^\top \bar{\mathbf{x}}_2 + \dots + \gamma^2 \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-2} + \gamma \boldsymbol{\theta}^\top \bar{\mathbf{x}}_{T-1} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}_T$$



Linear Recurrence

As T increases, we add new information x_T

Linear Recurrence

As T increases, we add new information x_T

As T increases, we slowly forget old information

Linear Recurrence

As T increases, we add new information x_T

As T increases, we slowly forget old information

The memory decay is smooth and differentiable

Linear Recurrence

As T increases, we add new information x_T

As T increases, we slowly forget old information

The memory decay is smooth and differentiable

We can learn the parameters γ, θ using gradient descent

Linear Recurrence

Morad et al., *Reinforcement Learning with Fast and Forgetful Memory*.
Neural Information Processing Systems. (2024).

Linear Recurrence




Morad et al., *Reinforcement Learning with Fast and Forgetful Memory*.
Neural Information Processing Systems. (2024).

$$\mathbf{H}_T = \gamma \mathbf{H}_{T-1} + g(\mathbf{x}_T)$$

Linear Recurrence

Morad et al., *Reinforcement Learning with Fast and Forgetful Memory*.
Neural Information Processing Systems. (2024).

$$\mathbf{H}_T = \gamma \mathbf{H}_{T-1} + g(\mathbf{x}_T)$$



<https://www.youtube.com/watch?v=0ey63XPB-4U&t=85s>

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. **Linear Recurrence**
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. **Scans**
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Scans

$$\textcolor{red}{h}_1 = f(\mathbf{0}, x_1, \theta) = \gamma \mathbf{0} + \theta^\top \bar{x}_1$$

$$\textcolor{green}{h}_2 = f(\textcolor{red}{h}_1, x_2, \theta) = \gamma h_1 + \theta^\top \bar{x}_2$$

$$h_3 = f(\textcolor{green}{h}_2, x_3, \theta) = \gamma h_2 + \theta^\top \bar{x}_3$$

\vdots

$$h_T = f(h_{T-1}, x_T, \theta) = \gamma h_{T-1} + \theta^\top \bar{x}_T$$

Scans

$$\textcolor{red}{h}_1 = f(\mathbf{0}, x_1, \theta) = \gamma \mathbf{0} + \theta^\top \bar{x}_1$$

$$\textcolor{green}{h}_2 = f(\textcolor{red}{h}_1, x_2, \theta) = \gamma \textcolor{red}{h}_1 + \theta^\top \bar{x}_2$$

$$h_3 = f(\textcolor{green}{h}_2, x_3, \theta) = \gamma h_2 + \theta^\top \bar{x}_3$$

\vdots

$$h_T = f(h_{T-1}, x_T, \theta) = \gamma h_{T-1} + \theta^\top \bar{x}_T$$

How do we compute h_1, h_2, \dots, h_T on a computer?

Scans

$$\textcolor{red}{h}_1 = f(\mathbf{0}, x_1, \theta) = \gamma \mathbf{0} + \theta^\top \bar{x}_1$$

$$\textcolor{green}{h}_2 = f(\textcolor{red}{h}_1, x_2, \theta) = \gamma \textcolor{red}{h}_1 + \theta^\top \bar{x}_2$$

$$h_3 = f(\textcolor{green}{h}_2, x_3, \theta) = \gamma h_2 + \theta^\top \bar{x}_3$$

\vdots

$$h_T = f(h_{T-1}, x_T, \theta) = \gamma h_{T-1} + \theta^\top \bar{x}_T$$

How do we compute h_1, h_2, \dots, h_T on a computer?

We use an algebraic operation called a **scan**

Scans

Our function f is just defined for a single X

Scans

Our function f is just defined for a single X

$$f : H \times X \times \Theta \mapsto H$$

Scans

Our function f is just defined for a single X

$$f : H \times X \times \Theta \mapsto H$$

To extend f to sequences, we scan f over the inputs

$$\text{scan}(f) : H \times X^T \times \Theta \mapsto H^T$$

Scans

Our function f is just defined for a single X

$$f : H \times X \times \Theta \mapsto H$$

To extend f to sequences, we scan f over the inputs

$$\text{scan}(f) : H \times X^T \times \Theta \mapsto H^T$$

$$\text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} h_1 \\ h_2 \\ \vdots \\ h_T \end{bmatrix}$$

Scans

$$f : H \times X \times \Theta \mapsto H, \quad \text{scan}(f) : H \times X^T \times \Theta \mapsto H^T$$

Scans

$$f : H \times X \times \Theta \mapsto H, \quad \text{scan}(f) : H \times X^T \times \Theta \mapsto H^T$$

$$\text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} h_1 \\ h_2 \\ \vdots \\ h_T \end{bmatrix}$$

Scans

$$f : H \times X \times \Theta \mapsto H, \quad \text{scan}(f) : H \times X^T \times \Theta \mapsto H^T$$

$$\text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} h_1 \\ h_2 \\ \vdots \\ h_T \end{bmatrix}$$

$$\text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} f(h_0, x_1, \theta) \\ f(h_1, x_2, \theta) \\ \vdots \\ f(h_{T-1}, x_T, \theta) \end{bmatrix} = \begin{bmatrix} f(h_0, x_1) \\ f(f(h_0, x_1), x_2) \\ \vdots \\ f(\dots f(h_0, x_1) \dots, x_T) \end{bmatrix}$$

Scans

```
import jax
import jax.numpy as jnp

T, d_x, d_h = 10, 2, 4
xs, h0 = jnp.ones((T, d_x)), jnp.zeros((d_h,))
theta = [jnp.ones((d_h,)), jnp.ones((d_x, d_h))] # (b, W)

def f(h, x):
    b, W = theta
    result = h + (W.T @ x + b)
    return result, result # return one, return all

hT, hs = jax.lax.scan(f, init=h0, xs=xs) # Scan f over x
```


Scans

torch does NOT have built-in scans, and is very slow compared to jax

Scans

torch does NOT have built-in scans, and is very slow compared to jax

We will write our own scan

Scans

torch does NOT have built-in scans, and is very slow compared to jax

We will write our own scan

```
def scan(f, h, xs):  
    # h shape is (d_h,)   
    # xs shape is (T, d_x)   
    hs = []  
    for x in xs:  
        h = f(h, x, theta)  
        hs.append(h)  
    # output shape is (T, d_h)   
    return torch.stack(hs)
```

Scans

```
import torch
T, d_x, d_h = 10, 2, 4

xs, h0 = torch.ones((T, d_x)), torch.zeros((d_h,))
theta = (torch.ones((d_h,)), torch.ones((d_x, d_h)))

def f(h, x):
    b, W = theta
    result = h + (W.T @ x + b)
    return result # h

hs = scan(f, h0, xs)
```

Scans

Many deep learning courses do not teach scans

Scans

Many deep learning courses do not teach scans

I teach scans because they are an important part of future LLMs

Scans

Many deep learning courses do not teach scans

I teach scans because they are an important part of future LLMs

OpenAI, Google, etc are currently experimenting with LLMs that use **associative scans**

Scans

Many deep learning courses do not teach scans

I teach scans because they are an important part of future LLMs

OpenAI, Google, etc are currently experimenting with LLMs that use **associative scans**

Standard LLM: ~8000 words, LLM + associative scan: 1M+ words

Scans

Many deep learning courses do not teach scans

I teach scans because they are an important part of future LLMs

OpenAI, Google, etc are currently experimenting with LLMs that use **associative scans**

Standard LLM: ~8000 words, LLM + associative scan: 1M+ words

An associative scan is a very fast scan we use when f obeys the associative property $(a + b) + c = a + (b + c)$

Scans

Many deep learning courses do not teach scans

I teach scans because they are an important part of future LLMs

OpenAI, Google, etc are currently experimenting with LLMs that use **associative scans**

Standard LLM: ~8000 words, LLM + associative scan: 1M+ words

An associative scan is a very fast scan we use when f obeys the associative property $(a + b) + c = a + (b + c)$

Question: Does $f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \overline{\mathbf{x}}$ obey the associative property?

Scans

Many deep learning courses do not teach scans

I teach scans because they are an important part of future LLMs

OpenAI, Google, etc are currently experimenting with LLMs that use **associative scans**

Standard LLM: ~8000 words, LLM + associative scan: 1M+ words

An associative scan is a very fast scan we use when f obeys the associative property $(a + b) + c = a + (b + c)$

Question: Does $f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}$ obey the associative property?

Answer: Yes, linear operations obey the associative property

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. **Scans**
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. **Output Modeling**
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Output Modeling

We are almost done defining recurrent models

Output Modeling

We are almost done defining recurrent models

There is one more step we must consider, **memory recall**

Output Modeling

We are almost done defining recurrent models

There is one more step we must consider, **memory recall**

h stores all memories, but humans only access a few memories at once

Output Modeling

We are almost done defining recurrent models

There is one more step we must consider, **memory recall**

h stores all memories, but humans only access a few memories at once

Example: I ask you your favorite ice cream flavor

Output Modeling

We are almost done defining recurrent models

There is one more step we must consider, **memory recall**

h stores all memories, but humans only access a few memories at once

Example: I ask you your favorite ice cream flavor

You recall previous times you ate ice cream, but not your phone number

Output Modeling

We are almost done defining recurrent models

There is one more step we must consider, **memory recall**

h stores all memories, but humans only access a few memories at once

Example: I ask you your favorite ice cream flavor

You recall previous times you ate ice cream, but not your phone number

We will model this recall of memories using a function g

Output Modeling

Let g define our memory recall function

Output Modeling

Let g define our memory recall function

$$g : H \times X \times \Theta \mapsto Y$$

Output Modeling

Let g define our memory recall function

$$g : H \times X \times \Theta \mapsto Y$$

g searches your memories h using the input x , to produce output y

Output Modeling

Let g define our memory recall function

$$g : H \times X \times \Theta \mapsto Y$$

g searches your memories h using the input x , to produce output y

x : “What is your favorite ice cream flavor?”

Output Modeling

Let g define our memory recall function

$$g : H \times X \times \Theta \mapsto Y$$

g searches your memories h using the input x , to produce output y

x : “What is your favorite ice cream flavor?”

h : Everything you remember (hometown, birthday, etc)

Output Modeling

Let g define our memory recall function

$$g : H \times X \times \Theta \mapsto Y$$

g searches your memories h using the input x , to produce output y

x : “What is your favorite ice cream flavor?”

h : Everything you remember (hometown, birthday, etc)

g : Searches your memories for ice cream information, and responds “chocolate”

Output Modeling

Let g define our memory recall function

$$g : H \times X \times \Theta \mapsto Y$$

g searches your memories h using the input x , to produce output y

x : “What is your favorite ice cream flavor?”

h : Everything you remember (hometown, birthday, etc)

g : Searches your memories for ice cream information, and responds “chocolate”

Now, we will combine f and g

Output Modeling

Step 1: Perform scan to find recurrent states

Output Modeling

Step 1: Perform scan to find recurrent states

$$\begin{bmatrix} \mathbf{h}_1 \\ \vdots \\ \mathbf{h}_T \end{bmatrix} = \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta}_f \right)$$

Output Modeling

Step 1: Perform scan to find recurrent states

$$\begin{bmatrix} \mathbf{h}_1 \\ \vdots \\ \mathbf{h}_T \end{bmatrix} = \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta}_f \right)$$

Step 2: Perform recall on recurrent states

Output Modeling

Step 1: Perform scan to find recurrent states

$$\begin{bmatrix} \mathbf{h}_1 \\ \vdots \\ \mathbf{h}_T \end{bmatrix} = \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta}_f \right)$$

Step 2: Perform recall on recurrent states

$$\begin{bmatrix} \mathbf{y}_1 \\ \vdots \\ \mathbf{y}_T \end{bmatrix} = \begin{bmatrix} g(\mathbf{h}_1, \mathbf{x}_1, \boldsymbol{\theta}_g) \\ \vdots \\ g(\mathbf{h}_T, \mathbf{x}_T, \boldsymbol{\theta}_g) \end{bmatrix}$$

Output Modeling

Step 1: Perform scan to find recurrent states

$$\begin{bmatrix} \mathbf{h}_1 \\ \vdots \\ \mathbf{h}_T \end{bmatrix} = \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta}_f \right)$$

Step 2: Perform recall on recurrent states

$$\begin{bmatrix} \mathbf{y}_1 \\ \vdots \\ \mathbf{y}_T \end{bmatrix} = \begin{bmatrix} g(\mathbf{h}_1, \mathbf{x}_1, \boldsymbol{\theta}_g) \\ \vdots \\ g(\mathbf{h}_T, \mathbf{x}_T, \boldsymbol{\theta}_g) \end{bmatrix}$$

Questions? This is on the homework

Output Modeling

To summarize, we defined:

Output Modeling

To summarize, we defined:

- Recurrent function f

Output Modeling

To summarize, we defined:

- Recurrent function f
- Scanned recurrence $\text{scan}(f)$

Output Modeling

To summarize, we defined:

- Recurrent function f
- Scanned recurrence $\text{scan}(f)$
- Output function g

To run our model:

Output Modeling

To summarize, we defined:

- Recurrent function f
- Scanned recurrence $\text{scan}(f)$
- Output function g

To run our model:

- Execute $\text{scan}(f)$ over inputs to make recurrent states

Output Modeling

To summarize, we defined:

- Recurrent function f
- Scanned recurrence $\text{scan}(f)$
- Output function g

To run our model:

- Execute $\text{scan}(f)$ over inputs to make recurrent states
- Execute g over recurrent states to make outputs

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. **Output Modeling**
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Relax

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. **Recurrent Loss Functions**
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Recurrent Loss Functions

Let us examine some example tasks:

Recurrent Loss Functions

Let us examine some example tasks:

- Clock

Recurrent Loss Functions

Let us examine some example tasks:

- Clock
- Explaining a video

Recurrent Loss Functions

Task: Clock – keep track of time

Recurrent Loss Functions

Task: Clock – keep track of time

Every minute, the minute hand ticks

Recurrent Loss Functions

Task: Clock – keep track of time

Every minute, the minute hand ticks

Count/remember the ticks to know the time

Recurrent Loss Functions

Task: Clock – keep track of time

Every minute, the minute hand ticks

Count/remember the ticks to know the time

$$X \in \{0, 1\}, \quad Y \in \mathbb{R}^2$$

Recurrent Loss Functions

Example input sequence:

$$\begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}$$

Recurrent Loss Functions

Example input sequence:

$$\begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix}$$

Desired output sequence

$$\begin{bmatrix} 0 & 1 \\ 0 & 2 \\ \vdots & \vdots \\ 2 & 13 \end{bmatrix}$$

We have a corresponding label y for each input x

Recurrent Loss Functions

Can use square error

Recurrent Loss Functions

Can use square error

First, scan f over the inputs to find h

$$h_{[i],j} = \text{scan}(f) \left(h_0, \begin{bmatrix} \mathbf{x}_{[i],1} \\ \vdots \\ \mathbf{x}_{[i],T} \end{bmatrix}, \theta_f \right)$$

Recurrent Loss Functions

Can use square error

First, scan f over the inputs to find h

$$h_{[i],j} = \text{scan}(f) \left(h_0, \begin{bmatrix} \mathbf{x}_{[i],1} \\ \vdots \\ \mathbf{x}_{[i],T} \end{bmatrix}, \theta_f \right)$$

$$\mathcal{L}(\mathbf{X}, \mathbf{Y}, \theta) = \sum_{i=1}^n \sum_{j=1}^T \left[g(h_{[i],j}, \mathbf{x}_{[i],j}, \theta_g) - \mathbf{y}_{[i],j} \right]^2$$

Recurrent Loss Functions

Can use square error

First, scan f over the inputs to find h

$$h_{[i],j} = \text{scan}(f) \left(h_0, \begin{bmatrix} \mathbf{x}_{[i],1} \\ \vdots \\ \mathbf{x}_{[i],T} \end{bmatrix}, \theta_f \right)$$

$$\mathcal{L}(\mathbf{X}, \mathbf{Y}, \theta) = \sum_{i=1}^n \sum_{j=1}^T \left[g(h_{[i],j}, \mathbf{x}_{[i],j}, \theta_g) - \mathbf{y}_{[i],j} \right]^2$$

Onto the next task

Recurrent Loss Functions

Task: Watch a video, then explain it

$$X \in \mathbb{Z}^{3 \times 32 \times 32}, \quad Y \in \{\text{comedy show, action movie, ...}\}$$

Recurrent Loss Functions

Task: Watch a video, then explain it

$$X \in \mathbb{Z}^{3 \times 32 \times 32}, \quad Y \in \{\text{comedy show, action movie, ...}\}$$

Example input sequence:

$$\begin{bmatrix} I_1 \\ I_2 \\ \vdots \\ I_T \end{bmatrix}$$

Example output:

“dancing dog”

Recurrent Loss Functions

Task: Watch a video, then explain it

$$X \in \mathbb{Z}^{3 \times 32 \times 32}, \quad Y \in \{\text{comedy show, action movie, ...}\}$$

Example input sequence:

$$\begin{bmatrix} I_1 \\ I_2 \\ \vdots \\ I_T \end{bmatrix}$$

Example output:

“dancing dog”

Unlike before, we have many inputs but just one output!

Recurrent Loss Functions

We will use the classification loss

Recurrent Loss Functions

We will use the classification loss

We scan f over the sequence, then compute g for the final timestep

Recurrent Loss Functions

We will use the classification loss

We scan f over the sequence, then compute g for the final timestep

$$h_{[i],j} = \text{scan}(f) \left(h_0, \begin{bmatrix} \mathbf{x}_{[i],1} \\ \vdots \\ \mathbf{x}_{[i],T} \end{bmatrix}, \theta_f \right)$$

Recurrent Loss Functions

We will use the classification loss

We scan f over the sequence, then compute g for the final timestep

$$\mathbf{h}_{[i],j} = \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_{[i],1} \\ \vdots \\ \mathbf{x}_{[i],T} \end{bmatrix}, \boldsymbol{\theta}_f \right)$$

$$\mathcal{L}(\mathbf{X}, \mathbf{Y}, \boldsymbol{\theta}) = \sum_{i=1}^n \sum_{j=1}^{d_y} y_{[i],j} \log g(\mathbf{h}_{[i],T}, \mathbf{x}_{[i],T}, \boldsymbol{\theta}_g)_j$$

Recurrent Loss Functions

We will use the classification loss

We scan f over the sequence, then compute g for the final timestep

$$\mathbf{h}_{[i],j} = \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_{[i],1} \\ \vdots \\ \mathbf{x}_{[i],T} \end{bmatrix}, \boldsymbol{\theta}_f \right)$$

$$\mathcal{L}(\mathbf{X}, \mathbf{Y}, \boldsymbol{\theta}) = \sum_{i=1}^n \sum_{j=1}^{d_y} y_{[i],j} \log g(\mathbf{h}_{[i],T}, \mathbf{x}_{[i],T}, \boldsymbol{\theta}_g)_j$$

We only care about the \mathbf{h}_T

Recurrent Loss Functions

To summarize, we use standard losses for recurrent loss functions

Recurrent Loss Functions

To summarize, we use standard losses for recurrent loss functions

Just be careful – we often sum over an additional axis

$$\sum_{i=1}^n \sum_{j=1}^T \dots$$

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. **Recurrent Loss Functions**
8. Backpropagation through Time
9. Recurrent Neural Networks
10. Coding

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. **Backpropagation through Time**
9. Recurrent Neural Networks
10. Coding

Backpropagation through Time

1. We created the model

Backpropagation through Time

1. We created the model
2. We found the loss function

Backpropagation through Time

1. We created the model
2. We found the loss function
3. Now we need to find the parameters!

Backpropagation through Time

1. We created the model
2. We found the loss function
3. Now we need to find the parameters!

Just like all other models, we train recurrent models using gradient descent

Backpropagation through Time

1. We created the model
2. We found the loss function
3. Now we need to find the parameters!

Just like all other models, we train recurrent models using gradient descent

Step 1: Compute gradient

Backpropagation through Time

1. We created the model
2. We found the loss function
3. Now we need to find the parameters!

Just like all other models, we train recurrent models using gradient descent

Step 1: Compute gradient

Step 2: Update parameters using gradient

Backpropagation through Time

1. We created the model
2. We found the loss function
3. Now we need to find the parameters!

Just like all other models, we train recurrent models using gradient descent

Step 1: Compute gradient

Step 2: Update parameters using gradient

How do we compute gradients for recurrent functions?

Backpropagation through Time

First, compute gradient of f

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \overline{\mathbf{x}}$$

Backpropagation through Time

First, compute gradient of f

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \overline{\mathbf{x}}$$

Question: What is $\nabla_{\boldsymbol{\theta}} f$?

Backpropagation through Time

First, compute gradient of f

$$f(h, x, \theta) = \gamma h + \theta^\top \bar{x}$$

Question: What is $\nabla_{\theta} f$?

$$\nabla_{\theta} f(h, x, \theta) = \bar{x}^\top$$

Backpropagation through Time

First, compute gradient of f

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \gamma \mathbf{h} + \boldsymbol{\theta}^\top \bar{\mathbf{x}}$$

Question: What is $\nabla_{\boldsymbol{\theta}} f$?

$$\nabla_{\boldsymbol{\theta}} f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \bar{\mathbf{x}}^\top$$

Too easy, now let us find the gradient of $\text{scan}(f)$

Backpropagation through Time

$$\text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} f(\mathbf{h}_0, \mathbf{x}_1, \boldsymbol{\theta}) \\ f(\mathbf{h}_1, \mathbf{x}_2, \boldsymbol{\theta}) \\ \vdots \\ f(\mathbf{h}_{T-1}, \mathbf{x}_T, \boldsymbol{\theta}) \end{bmatrix} = \begin{bmatrix} f(\mathbf{h}_0, \mathbf{x}_1) \\ f(f(\mathbf{h}_0, \mathbf{x}_1), \mathbf{x}_2) \\ \vdots \\ f(\dots f(\mathbf{h}_0, \mathbf{x}_1) \dots, \mathbf{x}_T) \end{bmatrix}$$

Question: What is $\nabla_{\boldsymbol{\theta}} \text{scan}(f)$?

Hint: Chain rule

Backpropagation through Time

$$\text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} f(\mathbf{h}_0, \mathbf{x}_1, \boldsymbol{\theta}) \\ f(\mathbf{h}_1, \mathbf{x}_2, \boldsymbol{\theta}) \\ \vdots \\ f(\mathbf{h}_{T-1}, \mathbf{x}_T, \boldsymbol{\theta}) \end{bmatrix} = \begin{bmatrix} f(\mathbf{h}_0, \mathbf{x}_1) \\ f(f(\mathbf{h}_0, \mathbf{x}_1), \mathbf{x}_2) \\ \vdots \\ f(\dots f(\mathbf{h}_0, \mathbf{x}_1) \dots, \mathbf{x}_T) \end{bmatrix}$$

Question: What is $\nabla_{\boldsymbol{\theta}} \text{scan}(f)$?

Hint: Chain rule

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f](\mathbf{h}_0, \mathbf{x}_1) \\ \nabla_{\boldsymbol{\theta}}[f](f(\mathbf{h}_0, \mathbf{x}_1), \mathbf{x}_2) \nabla_{\boldsymbol{\theta}}[f](\mathbf{h}_0, \mathbf{x}_1) \\ \vdots \\ \nabla_{\boldsymbol{\theta}}[f](\dots f(\mathbf{h}_0, \mathbf{x}_1) \dots, \mathbf{x}_T) \dots \cdot \nabla_{\boldsymbol{\theta}}[f](\mathbf{h}_0, \mathbf{x}_1) \end{bmatrix}$$

Backpropagation through Time

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\boldsymbol{h}_0, \begin{bmatrix} \boldsymbol{x}_1 \\ \vdots \\ \boldsymbol{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \vdots \end{bmatrix}$$

Backpropagation through Time

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\boldsymbol{h}_0, \begin{bmatrix} \boldsymbol{x}_1 \\ \vdots \\ \boldsymbol{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \vdots \end{bmatrix}$$

Question: Any issues with this?

Backpropagation through Time

$$\nabla_{\theta} \text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} \nabla_{\theta}[f] \\ \nabla_{\theta}[f] \nabla_{\theta}[f] \\ \nabla_{\theta}[f] \nabla_{\theta}[f] \nabla_{\theta}[f] \\ \vdots \end{bmatrix}$$

Question: Any issues with this?

What if $\nabla_{\theta} f$ is $\ll 1$ or $\gg 1$?

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. **Backpropagation through Time**
9. Recurrent Neural Networks
10. Coding

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. **Recurrent Neural Networks**
10. Coding

Recurrent Neural Networks

Until now, f was a linear function

Recurrent Neural Networks

Until now, f was a linear function

If we make f a neural network, then we have a **recurrent neural network** (RNN)

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$g(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \boldsymbol{\theta}_3^\top \mathbf{h}$$

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$g(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \boldsymbol{\theta}_3^\top \mathbf{h}$$

\mathbf{h} grows large and causes exploding gradients, σ should be sigmoid!

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$g(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \boldsymbol{\theta}_3^\top \mathbf{h}$$

\mathbf{h} grows large and causes exploding gradients, σ should be sigmoid!

Question: Max value of $\sigma(\mathbf{h})$?

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$g(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \boldsymbol{\theta}_3^\top \mathbf{h}$$

\mathbf{h} grows large and causes exploding gradients, σ should be sigmoid!

Question: Max value of $\sigma(\mathbf{h})$? 1!

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$g(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \boldsymbol{\theta}_3^\top \mathbf{h}$$

\mathbf{h} grows large and causes exploding gradients, σ should be sigmoid!

Question: Max value of $\sigma(\mathbf{h})$? 1!

Question: Anything missing from our linear model?

Recurrent Neural Networks

The simplest recurrent neural network is the **Elman Network**

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$g(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \boldsymbol{\theta}_3^\top \mathbf{h}$$

\mathbf{h} grows large and causes exploding gradients, σ should be sigmoid!

Question: Max value of $\sigma(\mathbf{h})$? 1!

Question: Anything missing from our linear model?

Answer: Forgetting!

Recurrent Neural Networks

Add forgetting

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} \odot f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

Recurrent Neural Networks

Add forgetting

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} \odot f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) + \boldsymbol{\theta}_2^\top \bar{\mathbf{x}})$$

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \bar{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

Recurrent Neural Networks

Add forgetting

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} \odot f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) + \boldsymbol{\theta}_2^\top \bar{\mathbf{x}})$$

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \bar{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

Question: σ is sigmoid. What is range/codomain of f_{forget} ?

Recurrent Neural Networks

Add forgetting

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} \odot f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) + \boldsymbol{\theta}_2^\top \bar{\mathbf{x}})$$

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \bar{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

Question: σ is sigmoid. What is range/codomain of f_{forget} ?

Answer: $[0, 1]$

Recurrent Neural Networks

Add forgetting

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} \odot f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) + \boldsymbol{\theta}_2^\top \bar{\mathbf{x}})$$

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \bar{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

Question: σ is sigmoid. What is range/codomain of f_{forget} ?

Answer: $[0, 1]$

When $f_{\text{forget}} < 1$, we forget some of our memories!

Recurrent Neural Networks

Add forgetting

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} \odot f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) + \boldsymbol{\theta}_2^\top \bar{\mathbf{x}})$$

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \bar{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

Question: σ is sigmoid. What is range/codomain of f_{forget} ?

Answer: $[0, 1]$

When $f_{\text{forget}} < 1$, we forget some of our memories!

Through gradient descent, neural network learns which memories to forget

Recurrent Neural Networks

Minimal gated unit (MGU) is a modern RNN

Recurrent Neural Networks

Minimal gated unit (MGU) is a modern RNN

MGU defines two helper functions

Recurrent Neural Networks

Minimal gated unit (MGU) is a modern RNN

MGU defines two helper functions

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \overline{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

Recurrent Neural Networks

Minimal gated unit (MGU) is a modern RNN

MGU defines two helper functions

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \overline{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

$$f_2(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_3^\top \overline{\mathbf{x}} + \boldsymbol{\theta}_4^\top f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) \odot \mathbf{h})$$

Recurrent Neural Networks

Minimal gated unit (MGU) is a modern RNN

MGU defines two helper functions

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \overline{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

$$f_2(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_3^\top \overline{\mathbf{x}} + \boldsymbol{\theta}_4^\top f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) \odot \mathbf{h})$$

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) \odot \mathbf{h} + (1 - f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta})) \odot f_2(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta})$$

Recurrent Neural Networks

Minimal gated unit (MGU) is a modern RNN

MGU defines two helper functions

$$f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \overline{\mathbf{x}} + \boldsymbol{\theta}_2^\top \mathbf{h})$$

$$f_2(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_3^\top \overline{\mathbf{x}} + \boldsymbol{\theta}_4^\top f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) \odot \mathbf{h})$$

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) \odot \mathbf{h} + (1 - f_{\text{forget}}(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta})) \odot f_2(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta})$$

Left term forgets old, right term replaces forgotten memories

Recurrent Neural Networks

There are even more complicated models

Recurrent Neural Networks

There are even more complicated models

- Long Short-Term Memory (LSTM)

Recurrent Neural Networks

There are even more complicated models

- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

Recurrent Neural Networks

There are even more complicated models

- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

LSTM has 6 different functions! Too complicated to review

Recurrent Neural Networks

There are even more complicated models

- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

LSTM has 6 different functions! Too complicated to review

MGU is simpler and performs similarly to LSTM and GRU

Recurrent Neural Networks

Recall the gradient for the linear recurrence

Recurrent Neural Networks

Recall the gradient for the linear recurrence

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\boldsymbol{h}_0, \begin{bmatrix} \boldsymbol{x}_1 \\ \vdots \\ \boldsymbol{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \vdots \end{bmatrix}$$

Recurrent Neural Networks

Recall the gradient for the linear recurrence

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \vdots \end{bmatrix}$$

Elman network f : $f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$

Recurrent Neural Networks

Recall the gradient for the linear recurrence

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \vdots \end{bmatrix}$$

Elman network f : $f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$

Question: What is the gradient for $\text{scan}(f)$ of Elman network?

Recurrent Neural Networks

Elman network f :

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

Recurrent Neural Networks

Elman network f :

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \vdots \end{bmatrix}$$

Recurrent Neural Networks

Elman network f :

$$f(\mathbf{h}, \mathbf{x}, \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta}_1^\top \mathbf{h} + \boldsymbol{\theta}_2^\top \overline{\mathbf{x}})$$

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \nabla_{\boldsymbol{\theta}}[f] \\ \vdots \end{bmatrix}$$

$$\nabla_{\boldsymbol{\theta}} \text{scan}(f) \left(\mathbf{h}_0, \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \boldsymbol{\theta} \right) = \begin{bmatrix} \nabla_{\boldsymbol{\theta}}[\sigma] \\ \nabla_{\boldsymbol{\theta}}[\sigma] \nabla_{\boldsymbol{\theta}}[\sigma] \\ \nabla_{\boldsymbol{\theta}}[\sigma] \nabla_{\boldsymbol{\theta}}[\sigma] \nabla_{\boldsymbol{\theta}}[\sigma] \\ \vdots \end{bmatrix}$$

Recurrent Neural Networks

$$\nabla_{\theta} \text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \vdots \end{bmatrix}$$

Question: What's the problem?

Recurrent Neural Networks

$$\nabla_{\theta} \text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \vdots \end{bmatrix}$$

Question: What's the problem?

Answer: Vanishing gradient

$$\nabla[\sigma] \cdot \nabla[\sigma] \cdot \dots = 0$$

Recurrent Neural Networks

$$\nabla_{\theta} \text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \vdots \end{bmatrix}$$

Question: What's the problem?

Answer: Vanishing gradient

$$\nabla[\sigma] \cdot \nabla[\sigma] \cdot \dots = 0$$

Question: What can we do?

Recurrent Neural Networks

$$\nabla_{\theta} \text{scan}(f) \left(h_0, \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta \right) = \begin{bmatrix} \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \nabla_{\theta}[\sigma] \\ \vdots \end{bmatrix}$$

Question: What's the problem?

Answer: Vanishing gradient

$$\nabla[\sigma] \cdot \nabla[\sigma] \cdot \dots = 0$$

Question: What can we do?

All RNNs suffer from either exploding gradient (ReLU) or vanishing gradient (sigmoid). Active area of research!

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. **Recurrent Neural Networks**
10. Coding

Agenda

1. Review
2. Sequence Modeling
3. Composite Memory
4. Linear Recurrence
5. Scans
6. Output Modeling
7. Recurrent Loss Functions
8. Backpropagation through Time
9. Recurrent Neural Networks
10. **Coding**

Coding

Jax RNN https://colab.research.google.com/drive/147z7FNGyERV8oQ_4gZmxDVdeoNt0hKta#scrollTo=TUMonlJ1u8Va

Homework <https://colab.research.google.com/drive/1CNaDxx1yJ4-phyMvgbxECL8ydZYBGQGt?usp=sharing>