



# Transformers

CISC 7026 - Introduction to Deep Learning

Steven Morad

University of Macau

Review .....	2
Going Deeper .....	11
Transformers .....	29
Positional Encoding .....	33
Text Transformers .....	47
Image Transformers .....	48
Unsupervised Training .....	49
World Models .....	51
Course Evaluation .....	53

# Review

---

# Review

Last time, we derived various forms of **attention**

We started with composite memory

# Review

Last time, we derived various forms of **attention**

We started with composite memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

# Review

Last time, we derived various forms of **attention**

We started with composite memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

Given large enough  $T$ , we will eventually run out of storage space

# Review

Last time, we derived various forms of **attention**

We started with composite memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

Given large enough  $T$ , we will eventually run out of storage space

The sum is a **lossy** operation that can store a limited amount of information

# Review

Last time, we derived various forms of **attention**

We started with composite memory

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$

Given large enough  $T$ , we will eventually run out of storage space

The sum is a **lossy** operation that can store a limited amount of information

# Review

So we introduced a forgetting term  $\gamma$

# Review

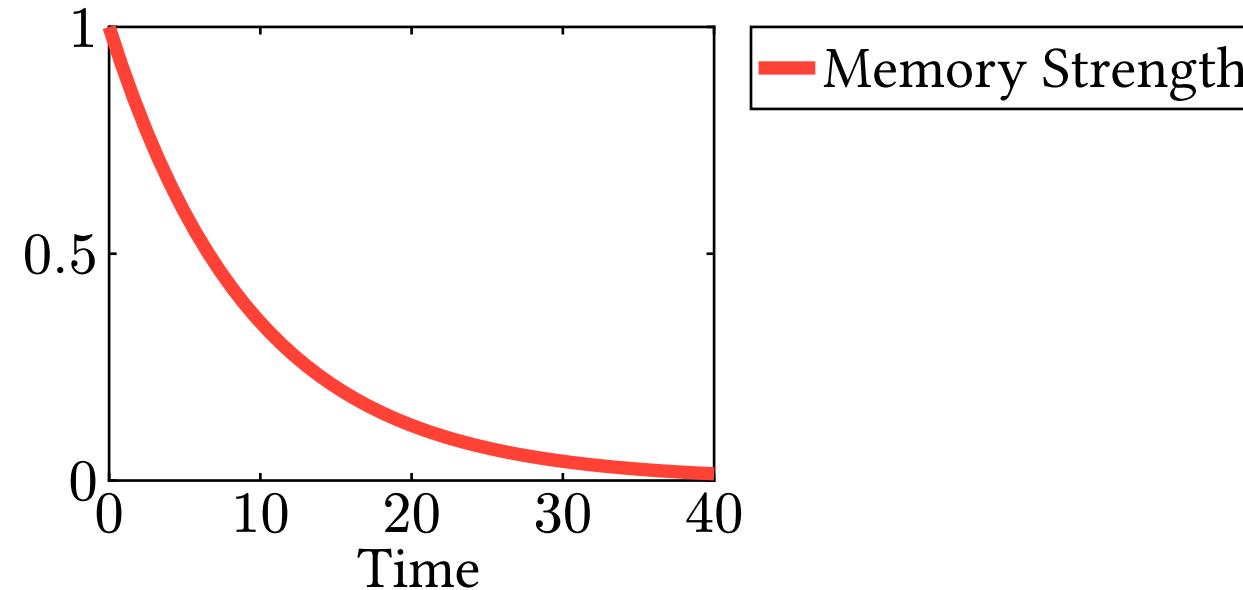
So we introduced a forgetting term  $\gamma$

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \gamma^{T-i} \cdot \boldsymbol{\theta}^\top \overline{\mathbf{x}}_i$$

# Review

So we introduced a forgetting term  $\gamma$

$$f(\mathbf{x}, \boldsymbol{\theta}) = \sum_{i=1}^T \gamma^{T-i} \cdot \boldsymbol{\theta}^\top \bar{\mathbf{x}}_i$$



# Review

We went to a party and the forgetting seemed ok

# Review

We went to a party and the forgetting seemed ok



# Review

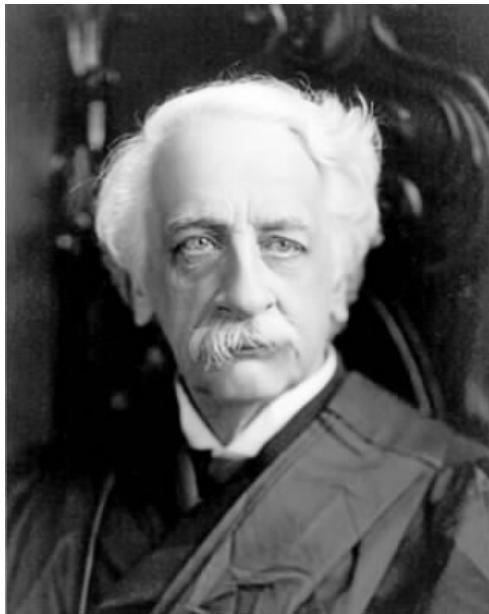
We went to a party and the forgetting seemed ok



10 PM

# Review

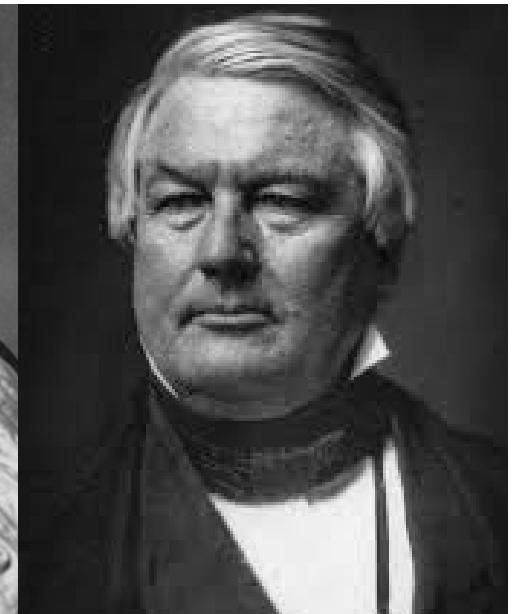
We went to a party and the forgetting seemed ok



10 PM

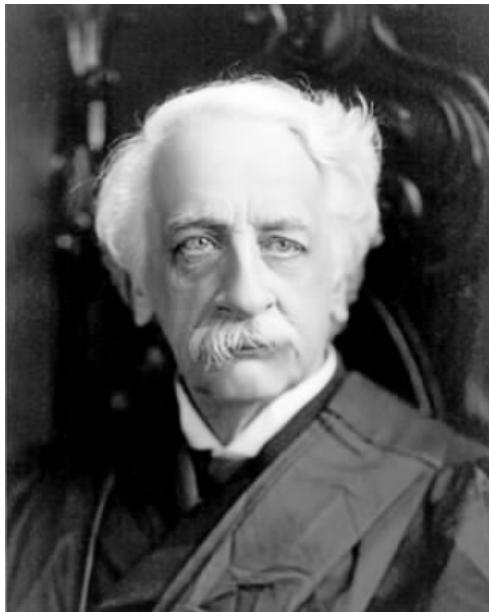


11 PM



# Review

We went to a party and the forgetting seemed ok



10 PM



11 PM

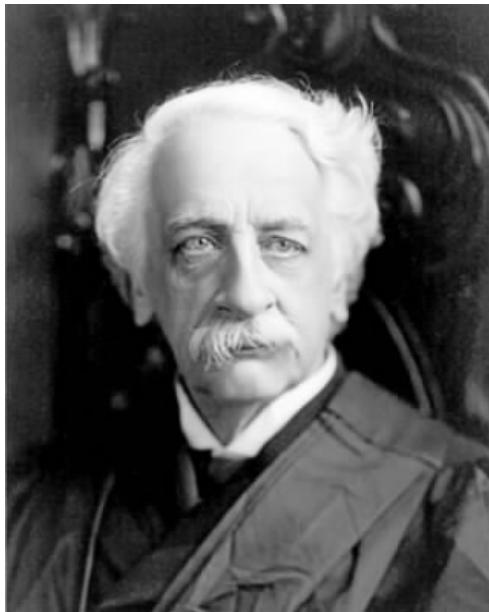


12 AM



# Review

We went to a party and the forgetting seemed ok



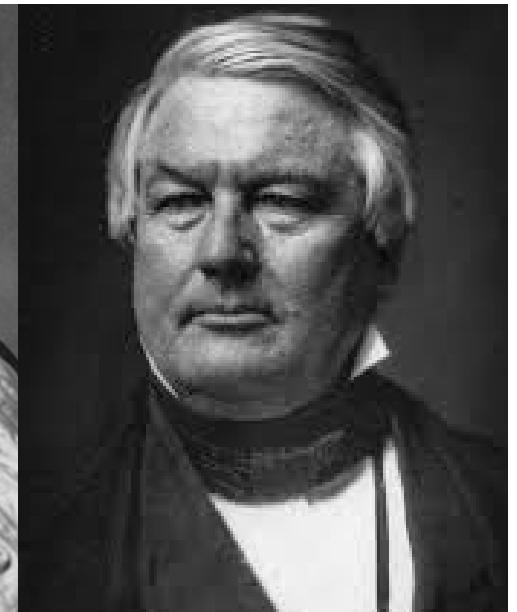
10 PM



11 PM



12 AM



1 AM

# Review



# Review



$$\gamma^3 \theta^\top \bar{x}_1$$

# Review



$$\gamma^3 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_1$$

$$\gamma^2 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_2$$

# Review



$$\gamma^3 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_1$$

$$\gamma^2 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_2$$

$$\gamma^1 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_3$$

# Review



$$\gamma^3 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_1$$

$$\gamma^2 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_2$$

$$\gamma^1 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_3$$

$$\gamma^0 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_4$$

# Review



$$\gamma^3 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_1$$

$$\gamma^2 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_2$$

$$\gamma^1 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_3$$

$$\gamma^0 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_4$$

# Review

But we encountered problems when Taylor Swift arrived at the party

# Review

But we encountered problems when Taylor Swift arrived at the party



# Review

But we encountered problems when Taylor Swift arrived at the party



$$\gamma^4 \theta^\top \bar{x}_1$$

$$\gamma^3 \theta^\top \bar{x}_2$$

$$\gamma^2 \theta^\top \bar{x}_3$$

$$\gamma^1 \theta^\top \bar{x}_4$$

$$\gamma^0 \theta^\top \bar{x}_5$$

# Review

But we encountered problems when Taylor Swift arrived at the party



$$\gamma^4 \theta^\top \bar{x}_1$$

$$\gamma^3 \theta^\top \bar{x}_2$$

$$\gamma^2 \theta^\top \bar{x}_3$$

$$\gamma^1 \theta^\top \bar{x}_4$$

$$\gamma^0 \theta^\top \bar{x}_5$$

# Review



$$\gamma^4 \theta^\top \bar{x}_1$$

$$\gamma^3 \theta^\top \bar{x}_2$$

$$\gamma^2 \theta^\top \bar{x}_3$$

$$\gamma^1 \theta^\top \bar{x}_4$$

$$\gamma^0 \theta^\top \bar{x}_5$$

# Review



$$\gamma^4 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_1$$

$$\gamma^3 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_2$$

$$\gamma^2 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_3$$

$$\gamma^1 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_4$$

$$\gamma^0 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_5$$

With our current model, we forget Taylor Swift!

# Review



$$\gamma^4 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_1$$

$$\gamma^3 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_2$$

$$\gamma^2 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_3$$

$$\gamma^1 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_4$$

$$\gamma^0 \boldsymbol{\theta}^\top \bar{\boldsymbol{x}}_5$$

With our current model, we forget Taylor Swift!

Our model of human memory is incomplete

# Review

# Review

Last time we studied attention

# Review

Last time we studied attention

Overview of transformer application and domains

# Going Deeper

---

# Going Deeper

We previously reviewed training tricks

# Going Deeper

We previously reviewed training tricks

- Deeper networks

# Going Deeper

We previously reviewed training tricks

- Deeper networks
- Parameter initialization

# Going Deeper

We previously reviewed training tricks

- Deeper networks
- Parameter initialization
- Stochastic gradient descent

# Going Deeper

We previously reviewed training tricks

- Deeper networks
- Parameter initialization
- Stochastic gradient descent
- Adaptive optimization

# Going Deeper

We previously reviewed training tricks

- Deeper networks
- Parameter initialization
- Stochastic gradient descent
- Adaptive optimization
- Weight decay

# Going Deeper

We previously reviewed training tricks

- Deeper networks
- Parameter initialization
- Stochastic gradient descent
- Adaptive optimization
- Weight decay

These methods empirically improve performance, but we do not always understand why

# Going Deeper

Modern transformers can be very deep

# Going Deeper

Modern transformers can be very deep

For this reason, they use two new training tricks to enable very deep models

# Going Deeper

Modern transformers can be very deep

For this reason, they use two new training tricks to enable very deep models

- Residual connections

# Going Deeper

Modern transformers can be very deep

For this reason, they use two new training tricks to enable very deep models

- Residual connections
- Layer normalization

# Going Deeper

Modern transformers can be very deep

For this reason, they use two new training tricks to enable very deep models

- Residual connections
- Layer normalization

Let us introduce these tricks

# Going Deeper

Modern transformers can be very deep

For this reason, they use two new training tricks to enable very deep models

- Residual connections
- Layer normalization

Let us introduce these tricks

We will start with the **residual connection**

# Going Deeper

Remember that a two-layer MLP is a universal function approximator

# Going Deeper

Remember that a two-layer MLP is a universal function approximator

$$| f(\mathbf{x}, \boldsymbol{\theta}) - g(\mathbf{x}) | < \varepsilon$$

# Going Deeper

Remember that a two-layer MLP is a universal function approximator

$$| f(\mathbf{x}, \boldsymbol{\theta}) - g(\mathbf{x}) | < \varepsilon$$

This is only as the width of the network goes to infinity

# Going Deeper

Remember that a two-layer MLP is a universal function approximator

$$| f(\mathbf{x}, \boldsymbol{\theta}) - g(\mathbf{x}) | < \varepsilon$$

This is only as the width of the network goes to infinity

For certain problems, we need deeper networks

# Going Deeper

Remember that a two-layer MLP is a universal function approximator

$$| f(\mathbf{x}, \boldsymbol{\theta}) - g(\mathbf{x}) | < \varepsilon$$

This is only as the width of the network goes to infinity

For certain problems, we need deeper networks

# Going Deeper

But there is a limit!

# Going Deeper

But there is a limit!

Making the network too deep can hurt performance

# Going Deeper

But there is a limit!

Making the network too deep can hurt performance

The theory is that the input information is **lost** somewhere in the network

# Going Deeper

But there is a limit!

Making the network too deep can hurt performance

The theory is that the input information is **lost** somewhere in the network

$$\mathbf{y} = f_k(\dots f_2(f_1(\mathbf{x}, \boldsymbol{\theta}_1), \boldsymbol{\theta}_2), \dots, \boldsymbol{\theta}_k)$$

# Going Deeper

$$\mathbf{y} = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

**Claim:** If the input information is present throughout the network, then we should be able to learn the identity function  $f(x) = x$

# Going Deeper

$$y = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

**Claim:** If the input information is present throughout the network, then we should be able to learn the identity function  $f(x) = x$

$$x = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

**Question:** We have seen a similar model, what was it?

# Going Deeper

$$y = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

**Claim:** If the input information is present throughout the network, then we should be able to learn the identity function  $f(x) = x$

$$x = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

**Question:** We have seen a similar model, what was it?

**Question:** Do you agree or disagree with the claim?

# Going Deeper

$$y = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

**Claim:** If the input information is present throughout the network, then we should be able to learn the identity function  $f(x) = x$

$$x = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

**Question:** We have seen a similar model, what was it?

**Question:** Do you agree or disagree with the claim?

<https://colab.research.google.com/drive/1qVlbQKpTuBYIa7FvC4IH-kJq-E0jmc0d#scrollTo=bg74S-AvbmJz>

# Going Deeper

$$x = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

# Going Deeper

$$x = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

Very deep networks struggle to learn the identity function

# Going Deeper

$$x = f_k(\dots f_2(f_1(x, \theta_1), \theta_2), \dots, \theta_k)$$

Very deep networks struggle to learn the identity function

If the input information is available, then learning the identity function should be very easy!

**Question:** How can we prevent the input from getting lost?

# Going Deeper

We can feed the input to each layer

# Going Deeper

We can feed the input to each layer

The first approach is called the **DenseNet** approach

# Going Deeper

We can feed the input to each layer

The first approach is called the **DenseNet** approach

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2\left(\begin{bmatrix} x \\ z_1 \end{bmatrix}, \theta_2\right)$$

⋮

$$z_k = f_k\left(\begin{bmatrix} x \\ z_1 \\ \vdots \\ z_{k-1} \end{bmatrix}, \theta_k\right)$$

# Going Deeper

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2\left(\begin{bmatrix} x \\ z_1 \end{bmatrix}, \theta_2\right)$$

⋮

$$z_k = f_k\left(\begin{bmatrix} x \\ z_1 \\ \vdots \\ z_{k-1} \end{bmatrix}, \theta_k\right)$$

# Going Deeper

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2\left(\begin{bmatrix} x \\ z_1 \end{bmatrix}, \theta_2\right)$$

⋮

$$z_k = f_k\left(\begin{bmatrix} x \\ z_1 \\ \vdots \\ z_{k-1} \end{bmatrix}, \theta_k\right)$$

**Question:** Any issues with the DenseNet approach?

# Going Deeper

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2\left(\begin{bmatrix} x \\ z_1 \end{bmatrix}, \theta_2\right)$$

⋮

$$z_k = f_k\left(\begin{bmatrix} x \\ z_1 \\ \vdots \\ z_{k-1} \end{bmatrix}, \theta_k\right)$$

**Question:** Any issues with the DenseNet approach?

**Answer:** Very deep networks require too many parameters!

# Going Deeper

The next method is called the **ResNet** approach

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2(x, \theta_2) + z_1$$

$$z_3 = f_2(x, \theta_3) + z_2$$

⋮

$$z_k = f_k(x, \theta_k) + z_{k-1}$$

# Going Deeper

The next method is called the **ResNet** approach

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2(x, \theta_2) + z_1$$

$$z_3 = f_2(x, \theta_3) + z_2$$

⋮

$$z_k = f_k(x, \theta_k) + z_{k-1}$$

This allows information to flow around the layers

# Going Deeper

The next method is called the **ResNet** approach

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2(x, \theta_2) + z_1$$

$$z_3 = f_2(x, \theta_3) + z_2$$

⋮

$$z_k = f_k(x, \theta_k) + z_{k-1}$$

This allows information to flow around the layers

It requires much fewer parameters than a dense net, but also does not work as well

# Going Deeper

The next method is called the **ResNet** approach

$$z_1 = f_1(x, \theta_1)$$

$$z_2 = f_2(x, \theta_2) + z_1$$

$$z_3 = f_2(x, \theta_3) + z_2$$

⋮

$$z_k = f_k(x, \theta_k) + z_{k-1}$$

This allows information to flow around the layers

It requires much fewer parameters than a dense net, but also does not work as well

# Going Deeper

We call  $f(x) + x$  a **residual connection**

# Going Deeper

We call  $f(x) + x$  a **residual connection**

Instead of learning how to change  $x$ ,  $f$  learns what to add to  $x$

# Going Deeper

We call  $f(\mathbf{x}) + \mathbf{x}$  a **residual connection**

Instead of learning how to change  $\mathbf{x}$ ,  $f$  learns what to add to  $\mathbf{x}$

For example, for an identity function we can easily learn

$$f(\mathbf{x}, \theta) = 0; \quad f(\mathbf{x}, \theta) + \mathbf{x} = \mathbf{x}$$

# Going Deeper

We call  $f(\mathbf{x}) + \mathbf{x}$  a **residual connection**

Instead of learning how to change  $\mathbf{x}$ ,  $f$  learns what to add to  $\mathbf{x}$

For example, for an identity function we can easily learn

$$f(\mathbf{x}, \theta) = 0; \quad f(\mathbf{x}, \theta) + \mathbf{x} = \mathbf{x}$$

This helps prevent information from getting lost in very deep networks

# Going Deeper

The second trick is called **layer normalization**

# Going Deeper

The second trick is called **layer normalization**

Recall that with parameter initialization and weight decay, we ensure the parameters are fairly small

# Going Deeper

The second trick is called **layer normalization**

Recall that with parameter initialization and weight decay, we ensure the parameters are fairly small

But we can still have very small or very large outputs from each layer

# Going Deeper

The second trick is called **layer normalization**

Recall that with parameter initialization and weight decay, we ensure the parameters are fairly small

But we can still have very small or very large outputs from each layer

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{d_x} \theta_{1,i} x_i$$

# Going Deeper

The second trick is called **layer normalization**

Recall that with parameter initialization and weight decay, we ensure the parameters are fairly small

But we can still have very small or very large outputs from each layer

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{d_x} \theta_{1,i} x_i$$

**Question:** If all  $x_i = 1$ ,  $\theta_{1,i} = 0.01$  and  $d_x = 1000$ , what is the output?

# Going Deeper

$$f_1(\boldsymbol{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{d_x} \theta_{1,i} x_i$$

# Going Deeper

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{d_x} \theta_{1,i} x_i$$

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{1000} 0.01 \cdot 1 = 10$$

# Going Deeper

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{d_x} \theta_{1,i} x_i$$

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{1000} 0.01 \cdot 1 = 10$$

What if we add another layer with the same  $d_x$  and  $\theta$ ?

$$f_2(\mathbf{z}, \boldsymbol{\theta}_2) = \sum_{i=1}^{1000} 0.01 \cdot 10 = 100$$

# Going Deeper

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{d_x} \theta_{1,i} x_i$$

$$f_1(\mathbf{x}, \boldsymbol{\theta}_1) = \sum_{i=1}^{1000} 0.01 \cdot 1 = 10$$

What if we add another layer with the same  $d_x$  and  $\theta$ ?

$$f_2(\mathbf{z}, \boldsymbol{\theta}_2) = \sum_{i=1}^{1000} 0.01 \cdot 10 = 100$$

**Question:** What is the problem?

# Going Deeper

Let us look at the gradient

# Going Deeper

Let us look at the gradient

$$\nabla_{\theta_1} f_2(f_1(x, \theta_1), \theta_2) =$$

# Going Deeper

Let us look at the gradient

$$\nabla_{\theta_1} f_2(f_1(x, \theta_1), \theta_2) = \nabla_{\theta_1}[f_2](f_1(x, \theta_1)) \cdot \nabla_{\theta_1}[f_1](x, \theta_1)$$

# Going Deeper

Let us look at the gradient

$$\begin{aligned}\nabla_{\theta_1} f_2(f_1(x, \theta_1), \theta_2) &= \nabla_{\theta_1}[f_2](f_1(x, \theta_1)) \cdot \nabla_{\theta_1}[f_1](x, \theta_1) \\ &\approx 100 \cdot 10\end{aligned}$$

Can cause exploding or vanishing gradient

# Going Deeper

Let us look at the gradient

$$\begin{aligned}\nabla_{\theta_1} f_2(f_1(x, \theta_1), \theta_2) &= \nabla_{\theta_1}[f_2](f_1(x, \theta_1)) \cdot \nabla_{\theta_1}[f_1](x, \theta_1) \\ &\approx 100 \cdot 10\end{aligned}$$

Can cause exploding or vanishing gradient

Deeper network  $\Rightarrow$  worse exploding/vanishing issues

**Question:** What can we do?

# Going Deeper

Let us look at the gradient

$$\begin{aligned}\nabla_{\theta_1} f_2(f_1(x, \theta_1), \theta_2) &= \nabla_{\theta_1}[f_2](f_1(x, \theta_1)) \cdot \nabla_{\theta_1}[f_1](x, \theta_1) \\ &\approx 100 \cdot 10\end{aligned}$$

Can cause exploding or vanishing gradient

Deeper network  $\Rightarrow$  worse exploding/vanishing issues

**Question:** What can we do?

# Going Deeper

We can use **layer normalization**

# Going Deeper

We can use **layer normalization**

First, layer normalization **centers** the output of the layer

# Going Deeper

We can use **layer normalization**

First, layer normalization **centers** the output of the layer

$$\mu = d_y \sum_{i=1}^{d_y} f(x, \theta)_i$$

$$f(x, \theta) - \mu$$

**Question:** What does this do?

# Going Deeper

We can use **layer normalization**

First, layer normalization **centers** the output of the layer

$$\mu = d_y \sum_{i=1}^{d_y} f(x, \theta)_i$$

$$f(x, \theta) - \mu$$

**Question:** What does this do?

**Answer:** Makes output have zero mean (both positive and negative values)

# Going Deeper

We can use **layer normalization**

First, layer normalization **centers** the output of the layer

$$\mu = d_y \sum_{i=1}^{d_y} f(x, \theta)_i$$

$$f(x, \theta) - \mu$$

**Question:** What does this do?

**Answer:** Makes output have zero mean (both positive and negative values)

# Going Deeper

$$\mu = d_y \sum_{i=1}^{d_y} f(\mathbf{x}, \boldsymbol{\theta})_i \quad f(\mathbf{x}, \boldsymbol{\theta}) - \mu$$

Then, layer normalization **rescales** the outputs

$$\sigma = \frac{\sqrt{\sum_{i=1}^{d_y} f(\mathbf{x}, \boldsymbol{\theta}_i - \mu)^2}}{d_y}$$

$$\text{LN}(f(\mathbf{x}, \boldsymbol{\theta})) = \frac{f(\mathbf{x}, \boldsymbol{\theta}) - \mu}{\sigma}$$

Now, the output of the layer is normally distributed

# Going Deeper

If the output is normally distributed:

# Going Deeper

If the output is normally distributed:

- 99.7% of outputs  $\in [-3, 3]$

# Going Deeper

If the output is normally distributed:

- 99.7% of outputs  $\in [-3, 3]$
- 99.99% of outputs  $\in [-4, 4]$

# Going Deeper

If the output is normally distributed:

- 99.7% of outputs  $\in [-3, 3]$
- 99.99% of outputs  $\in [-4, 4]$
- 99.9999% of outputs  $\in [-5, 5]$

# Going Deeper

If the output is normally distributed:

- 99.7% of outputs  $\in [-3, 3]$
- 99.99% of outputs  $\in [-4, 4]$
- 99.9999% of outputs  $\in [-5, 5]$

This helps prevent vanishing and exploding gradients

# Going Deeper

Now, let's combine residual connections and layer norm and try our very deep network again

TODO COLAB

# Transformers

---

# Transformers

Now we have everything we need to implement a transformer

# Transformers

Now we have everything we need to implement a transformer

- Attention

# Transformers

Now we have everything we need to implement a transformer

- Attention
- MLP

# Transformers

Now we have everything we need to implement a transformer

- Attention
- MLP
- Residual connections

# Transformers

Now we have everything we need to implement a transformer

- Attention
- MLP
- Residual connections
- Layer normalization

A deep neural network consists of many layers

# Transformers

Now we have everything we need to implement a transformer

- Attention
- MLP
- Residual connections
- Layer normalization

A deep neural network consists of many layers

A transformer consists of many **transformer blocks**

# Transformers

```
class TransformerBlock(nn.Module):
    def __init__(self):
        self.attn = Attention()
        self.mlp = Sequential(
            Linear(d_h, d_h), LeakyReLU(), Linear(d_h, d_h))
        self.norm1 = nn.LayerNorm(d_h)
        self.norm2 = nn.LayerNorm(d_h)

    def forward(self, x):
        x = self.norm1(self.attn(x) + x)
        x = self.norm2(self.mlp(x) + x)
        return x
```

# Transformers

```
class Transformer(nn.Module):
    def __init__(self):
        self.block1 = TransformerBlock()
        self.block2 = TransformerBlock()
        self.block3 = TransformerBlock()

    def forward(self, x):
        x = self.block1(x)
        x = self.block2(x)
        x = self.block3(x)
        return x
```

# Transformers

```
class Transformer(nn.Module):
    def __init__(self):
        self.block1 = TransformerBlock()
        self.block2 = TransformerBlock()
        self.block3 = TransformerBlock()

    def forward(self, x):
        x = self.block1(x)
        x = self.block2(x)
        x = self.block3(x)
        return x
```

**Question:** What are the input/output shapes of the transformer?

# Transformers

```
class Transformer(nn.Module):
    def __init__(self):
        self.block1 = TransformerBlock()
        self.block2 = TransformerBlock()
        self.block3 = TransformerBlock()

    def forward(self, x):
        x = self.block1(x)
        x = self.block2(x)
        x = self.block3(x)
        return x
```

**Question:** What are the input/output shapes of the transformer?

**Answer:**  $f : \mathbb{R}^{T \times d_x} \mapsto \mathbb{R}^{T \times d_h}$

# Positional Encoding

---

# Positional Encoding

Our transformer maps  $T$  inputs to  $T$  outputs

# Positional Encoding

Our transformer maps  $T$  inputs to  $T$  outputs

- $T$  words in a sentence

# Positional Encoding

Our transformer maps  $T$  inputs to  $T$  outputs

- $T$  words in a sentence
- $T$  pixels in an image

# Positional Encoding

Our transformer maps  $T$  inputs to  $T$  outputs

- $T$  words in a sentence
- $T$  pixels in an image
- $T$  amino acids in a protein

# Positional Encoding

Our transformer maps  $T$  inputs to  $T$  outputs

- $T$  words in a sentence
- $T$  pixels in an image
- $T$  amino acids in a protein

**Question:** Does the order of the  $T$  inputs matter?

# Positional Encoding

Our transformer maps  $T$  inputs to  $T$  outputs

- $T$  words in a sentence
- $T$  pixels in an image
- $T$  amino acids in a protein

**Question:** Does the order of the  $T$  inputs matter?

**Answer:** In some tasks, yes! In others, no

# Positional Encoding

**Question:** Detecting birds in an image of  $T$  pixels, does order matter?

# Positional Encoding

**Question:** Detecting birds in an image of  $T$  pixels, does order matter?



# Positional Encoding

**Question:** Detecting birds in an image of  $T$  pixels, does order matter?



**Answer:** Yes!

# Positional Encoding

**Question:**  $T$  robots searching for an object. Does order matter?

# Positional Encoding

**Question:**  $T$  robots searching for an object. Does order matter?



# Positional Encoding

**Question:**  $T$  robots searching for an object. Does order matter?



**Answer:** No!

# Positional Encoding

**Question:** We translate a sentence containing  $T$  words. Does order matter?

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{dog} \\ \text{licks} \\ \text{the} \\ \text{owner} \end{bmatrix}$$

# Positional Encoding

**Question:** We translate a sentence containing  $T$  words. Does order matter?

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{dog} \\ \text{licks} \\ \text{the} \\ \text{owner} \end{bmatrix}$$

$$\begin{bmatrix} x_1 \\ \textcolor{red}{x_5} \\ x_3 \\ x_4 \\ \textcolor{red}{x_2} \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{owner} \\ \text{licks} \\ \text{the} \\ \text{dog} \end{bmatrix}$$

# Positional Encoding

**Question:** We translate a sentence containing  $T$  words. Does order matter?

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{dog} \\ \text{licks} \\ \text{the} \\ \text{owner} \end{bmatrix}$$

$$\begin{bmatrix} x_1 \\ \textcolor{red}{x_5} \\ x_3 \\ x_4 \\ \textcolor{red}{x_2} \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{owner} \\ \text{licks} \\ \text{the} \\ \text{dog} \end{bmatrix}$$

**Answer:** Yes!

# Positional Encoding

For some tasks, we must know the order of the inputs

# Positional Encoding

For some tasks, we must know the order of the inputs

**Question:** Does the transformer know the order of the  $T$  inputs?

# Positional Encoding

For some tasks, we must know the order of the inputs

**Question:** Does the transformer know the order of the  $T$  inputs?

Define a **permutation matrix**  $P \in \{0, 1\}^{T \times T}$  that reorders the inputs

# Positional Encoding

For some tasks, we must know the order of the inputs

**Question:** Does the transformer know the order of the  $T$  inputs?

Define a **permutation matrix**  $P \in \{0, 1\}^{T \times T}$  that reorders the inputs

# Positional Encoding

Example 1:

$$P = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}; \quad a = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix};$$

# Positional Encoding

Example 1:

$$P = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}; \quad a = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}; \quad Pa = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}$$

# Positional Encoding

**Example 1:**

$$P = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}; \quad a = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}; \quad Pa = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}$$

**Example 2:**

$$P = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}; \quad a = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix};$$

# Positional Encoding

**Example 1:**

$$P = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}; \quad a = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}; \quad Pa = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}$$

**Example 2:**

$$P = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}; \quad a = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}; \quad Pa = \begin{bmatrix} 4 \\ 3 \\ 5 \end{bmatrix}$$

# Positional Encoding

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$$

# Positional Encoding

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$$

$f$  is equivariant, order **does** matter

# Positional Encoding

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$$

$f$  is equivariant, order **does** matter

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$$

# Positional Encoding

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$$

$f$  is equivariant, order **does** matter

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf \begin{pmatrix} x_1 \\ \vdots \\ x_n \end{pmatrix}$$

$f$  is equivariant, order **does not** matter

Which is a transformer?

# Positional Encoding

Recall dot product self attention

# Positional Encoding

Recall dot product self attention

$$Q = [q_1 \ q_2 \ \dots \ q_T] = [\theta_Q^\top x_1 \ \theta_Q^\top x_2 \ \dots \ \theta_Q^\top x_T]$$

$$K = [k_1 \ k_2 \ \dots \ k_T] = [\theta_K^\top x_1 \ \theta_K^\top x_2 \ \dots \ \theta_K^\top x_T]$$

$$V = [v_1 \ v_2 \ \dots \ v_T] = [\theta_V^\top x_1 \ \theta_V^\top x_2 \ \dots \ \theta_V^\top x_T]$$

# Positional Encoding

Recall dot product self attention

$$Q = [q_1 \ q_2 \ \dots \ q_T] = [\theta_Q^\top x_1 \ \theta_Q^\top x_2 \ \dots \ \theta_Q^\top x_T]$$

$$K = [k_1 \ k_2 \ \dots \ k_T] = [\theta_K^\top x_1 \ \theta_K^\top x_2 \ \dots \ \theta_K^\top x_T]$$

$$V = [v_1 \ v_2 \ \dots \ v_T] = [\theta_V^\top x_1 \ \theta_V^\top x_2 \ \dots \ \theta_V^\top x_T]$$

$$\text{attn}\left(\begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_h}}\right)V$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\frac{(\mathbf{PQ})(\mathbf{PK})^\top}{\sqrt{d_h}}\right)(\mathbf{PV})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\frac{(\mathbf{PQ})(\mathbf{PK})^\top}{\sqrt{d_h}}\right)(\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{PK})^\top}{\sqrt{d_h}}\right)(\mathbf{PV})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\frac{(\mathbf{PQ})(\mathbf{PK})^\top}{\sqrt{d_h}}\right)(\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{PK})^\top}{\sqrt{d_h}}\right)(\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{K}^\top \mathbf{P}^\top)}{\sqrt{d_h}}\right)(\mathbf{PV})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\frac{(\mathbf{PQ})(\mathbf{PK})^\top}{\sqrt{d_h}}\right)(\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{PK})^\top}{\sqrt{d_h}}\right)(\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{K}^\top \mathbf{P}^\top)}{\sqrt{d_h}}\right)(\mathbf{PV})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{K}^\top \mathbf{P}^\top)}{\sqrt{d_h}}\right)(\mathbf{P}\mathbf{V})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{K}^\top \mathbf{P}^\top)}{\sqrt{d_h}}\right)(\mathbf{P}\mathbf{V})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right)(\mathbf{P}\mathbf{V})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{K}^\top \mathbf{P}^\top)}{\sqrt{d_h}}\right)(\mathbf{P}\mathbf{V})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right)(\mathbf{P}\mathbf{V})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right)(\mathbf{P}\mathbf{V})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}(\mathbf{K}^\top \mathbf{P}^\top)}{\sqrt{d_h}}\right)(\mathbf{P}\mathbf{V})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \text{softmax}\left(\mathbf{P} \frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right)(\mathbf{P}\mathbf{V})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right)(\mathbf{P}\mathbf{V})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right) (\mathbf{PV})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right) (\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}}\right) \mathbf{P}^\top (\mathbf{PV})$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right) (\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}}\right) \mathbf{P}^\top (\mathbf{PV})$$

$$\mathbf{P}^\top \mathbf{P} = \mathbf{I};$$

# Positional Encoding

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}} \mathbf{P}^\top\right) (\mathbf{PV})$$

$$\text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}}\right) \mathbf{P}^\top (\mathbf{PV})$$

$$\mathbf{P}^\top \mathbf{P} = \mathbf{I}; \quad \text{attn}\left(\mathbf{P} \begin{bmatrix} \mathbf{x}_1 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}, \theta\right) = \mathbf{P} \text{softmax}\left(\frac{\mathbf{QK}^\top}{\sqrt{d_h}}\right) \mathbf{V}$$

# Positional Encoding

$$\text{attn}\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = P \text{softmax}\left(\frac{QK^\top}{\sqrt{d_h}}\right)V$$

# Positional Encoding

$$\text{attn}\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = P \text{softmax}\left(\frac{QK^\top}{\sqrt{d_h}}\right)V$$

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf\begin{pmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \end{pmatrix}$$

# Positional Encoding

$$\text{attn}\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = P \text{softmax}\left(\frac{QK^\top}{\sqrt{d_h}}\right)V$$

$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf\begin{pmatrix} \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \end{pmatrix}$$

**Question:** What does this mean?

# Positional Encoding

$$\text{attn}\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_T \end{bmatrix}, \theta\right) = P \text{softmax}\left(\frac{QK^\top}{\sqrt{d_h}}\right)V$$
$$f\left(P \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right) = Pf\left(\begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}\right)$$

**Question:** What does this mean?

**Answer:** Attention is equivariant, order **does not** matter

# Positional Encoding

Transformer cannot determine order of inputs! **Equivariant** to ordering

# Positional Encoding

Transformer cannot determine order of inputs! **Equivariant** to ordering

The following sentences are the same to a transformer

$$\begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \mathbf{x}_3 \\ \mathbf{x}_4 \\ \mathbf{x}_5 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{dog} \\ \text{licks} \\ \text{the} \\ \text{owner} \end{bmatrix}$$

# Positional Encoding

Transformer cannot determine order of inputs! **Equivariant** to ordering

The following sentences are the same to a transformer

$$\begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \mathbf{x}_3 \\ \mathbf{x}_4 \\ \mathbf{x}_5 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{dog} \\ \text{licks} \\ \text{the} \\ \text{owner} \end{bmatrix}$$

$$\begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_5 \\ \mathbf{x}_3 \\ \mathbf{x}_4 \\ \mathbf{x}_2 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{owner} \\ \text{licks} \\ \text{the} \\ \text{dog} \end{bmatrix}$$

# Positional Encoding

Transformer cannot determine order of inputs! **Equivariant** to ordering

The following sentences are the same to a transformer

$$\begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \mathbf{x}_3 \\ \mathbf{x}_4 \\ \mathbf{x}_5 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{dog} \\ \text{licks} \\ \text{the} \\ \text{owner} \end{bmatrix}$$

$$\begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_5 \\ \mathbf{x}_3 \\ \mathbf{x}_4 \\ \mathbf{x}_2 \end{bmatrix} = \begin{bmatrix} \text{The} \\ \text{owner} \\ \text{licks} \\ \text{the} \\ \text{dog} \end{bmatrix}$$

This is a problem! For some tasks, we care about input order

# Text Transformers

---

# Image Transformers

---

# Unsupervised Training

---

# Unsupervised Training

Predict the future

# World Models

---

# World Models

What if transformers could interact with the world?

# Course Evaluation

---

# Course Evaluation

Department instructed me to ask you for course feedback

# Course Evaluation

Department instructed me to ask you for course feedback

We take this feedback seriously

# Course Evaluation

Department instructed me to ask you for course feedback

We take this feedback seriously

Your feedback will impact future courses (and my job)

# Course Evaluation

Department instructed me to ask you for course feedback

We take this feedback seriously

Your feedback will impact future courses (and my job)

Be specific on what you like and do not like

# Course Evaluation

Department instructed me to ask you for course feedback

We take this feedback seriously

Your feedback will impact future courses (and my job)

Be specific on what you like and do not like

Your likes/dislikes will filter into your future courses

# Course Evaluation

Department instructed me to ask you for course feedback

We take this feedback seriously

Your feedback will impact future courses (and my job)

Be specific on what you like and do not like

Your likes/dislikes will filter into your future courses

If you are not comfortable writing English, write Chinese

# Course Evaluation

Department instructed me to ask you for course feedback

We take this feedback seriously

Your feedback will impact future courses (and my job)

Be specific on what you like and do not like

Your likes/dislikes will filter into your future courses

If you are not comfortable writing English, write Chinese

# Course Evaluation

I must leave the room to let you fill out this form

# Course Evaluation

I must leave the room to let you fill out this form

Please scan the QR code and complete the survey

# Course Evaluation

I must leave the room to let you fill out this form

Please scan the QR code and complete the survey

Department has suggested 10 minutes

# Course Evaluation

I must leave the room to let you fill out this form

Please scan the QR code and complete the survey

Department has suggested 10 minutes

<https://isw.um.edu.mo/siaweb>

# Course Evaluation

Research data labeling and collection

If you participated, come up